

Kubernetes 설계 개념

Source: 2018 Saad Ali, Google
Kubernetes Design Principles:
Understand the Why

개요

- Goal: A deeper understanding of Kubernetes
- Important tool for learning
 - Understand the problem
 - The “**why**” not just the “what”

Kubernetes


- Containerization was the key
 - Consistent, repeatable, reliable deployments on a **wide variety of systems.**
- Who will manage it?
 - You? Scripts? A system you write?
- **Kubernetes manages your cluster!**
 - **Deploys & monitors containerized workloads.**

How to deploy a workload?

- Obvious solution



Start container X on Node B



Node A



Node B

How to deploy a workload?


- Obvious solution
 - Problems with this approach?



Start containers X on Node B



Node A



Node B

- User has to

- Monitor and store state of every container/node
- "Catch up" any failed nodes that missed calls

- Complex, custom logic

+ What if:

- + Container crashes and dies?
- + What if node crashes and dies?
- + What if node B is momentarily not healthy?

Principle #1

Kubernetes APIs are declarative = 원하는 것을 서술하는 식


rather the imperative = 세부적인 명령 대신에

Declarative APIs

- Before:
 - **You:** provide **exact set of instructions** to drive to desired state
 - **System:** executes instructions
 - **You:** **monitor system**, and provide **further instructions** if it deviates.
- After:
 - **You:** define desired state
 - **System:** works to drive towards that state


How to deploy a workload?

- The Kubernetes way!
 - **You:** **create** API **object** that is persisted **on kube API server** until deletion
 - **System:** **all components work in parallel to drive to that state**




How to deploy a workload?

- The Kubernetes way!
 - **You:** create API object that is persisted on kube API server until deletion
 - **System:** all components work in parallel to drive to that state



Why declarative over imperative?

- Automatic recovery!
- Example:
 - Step 1: Node failure
 - Step 2: System automatically moves pod to healthy node




How to deploy a workload?

- How does node figure out what to do?
- Problems with this approach?



✦ What if:

- ✦ Container crashes and dies?
- ✦ What if node crashes and dies?
- ✦ What if node B is momentary not healthy?



- Master has to

- Monitor and store state of every component.
- "Catch up" any failed components that missed calls

Master becomes:

- Complex
- Fragile
- Difficult to extend

Principle #2

Principle #1 Kubernetes APIs are declarative rather than imperative

**Principle #2 The Kubernetes control plane is transparent.
There are no hidden internal APIs.**


즉, API server 뒷단에서 다시 Master가 대신 명령하지 않는다

No hidden internal APIs

- Before:
 - **Master:** provide exact set of instructions to drive to node to desired state
 - **Node:** executes instructions
 - **Master:** monitor nodes, and provide further instructions if state deviates.
- After:
 - **Master:** define desired state of node
 - **Node:** works independently to drive itself towards that state


No hidden internal APIs

- All components watch the Kubernetes API, and figure out what they need to do.



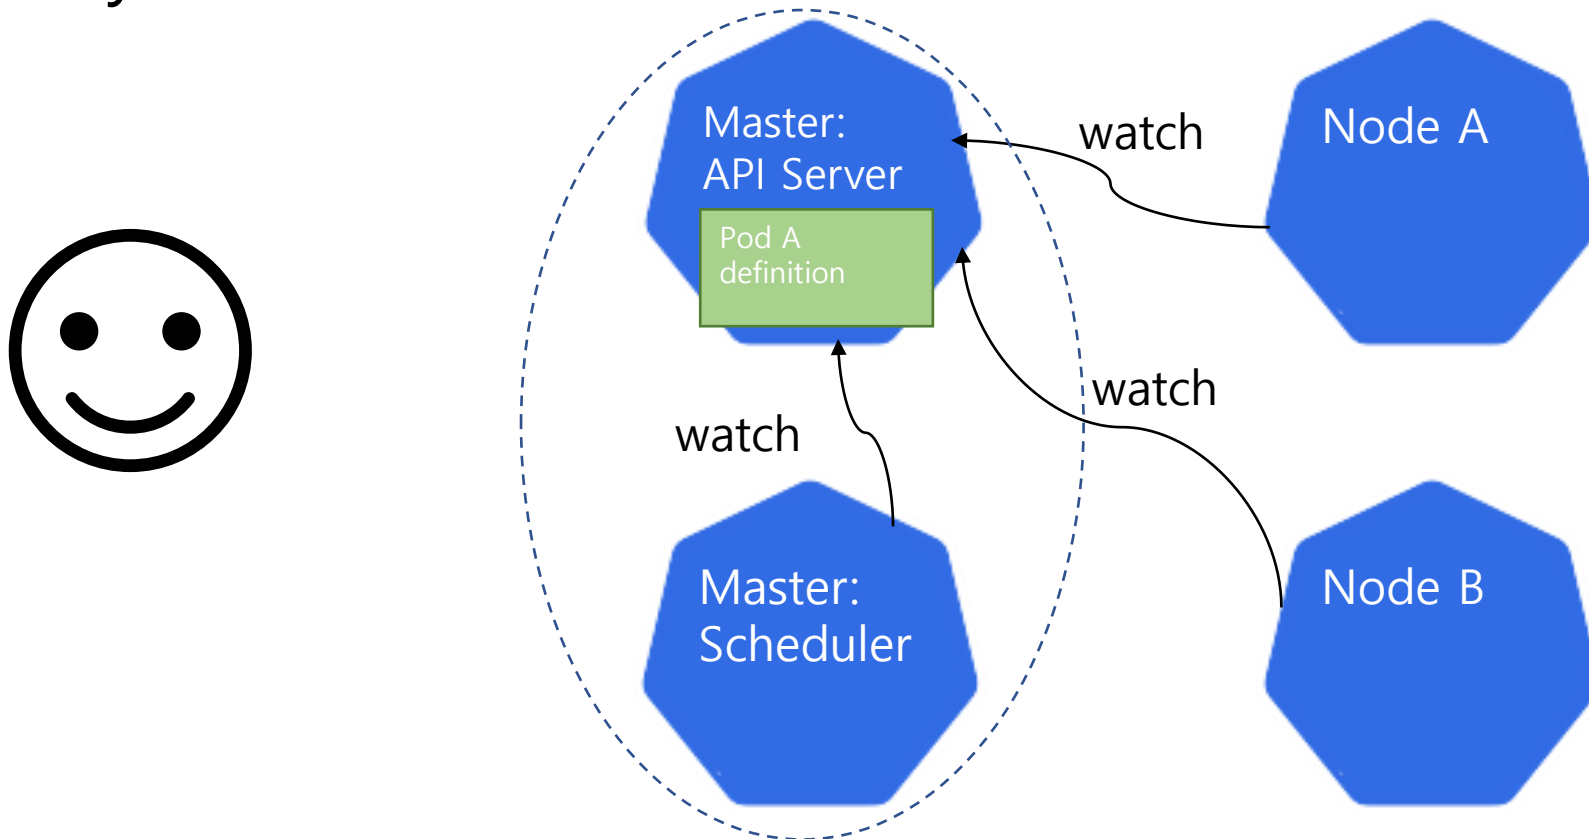
No hidden internal APIs

- All components watch the Kubernetes API, and figure out what they need to do.



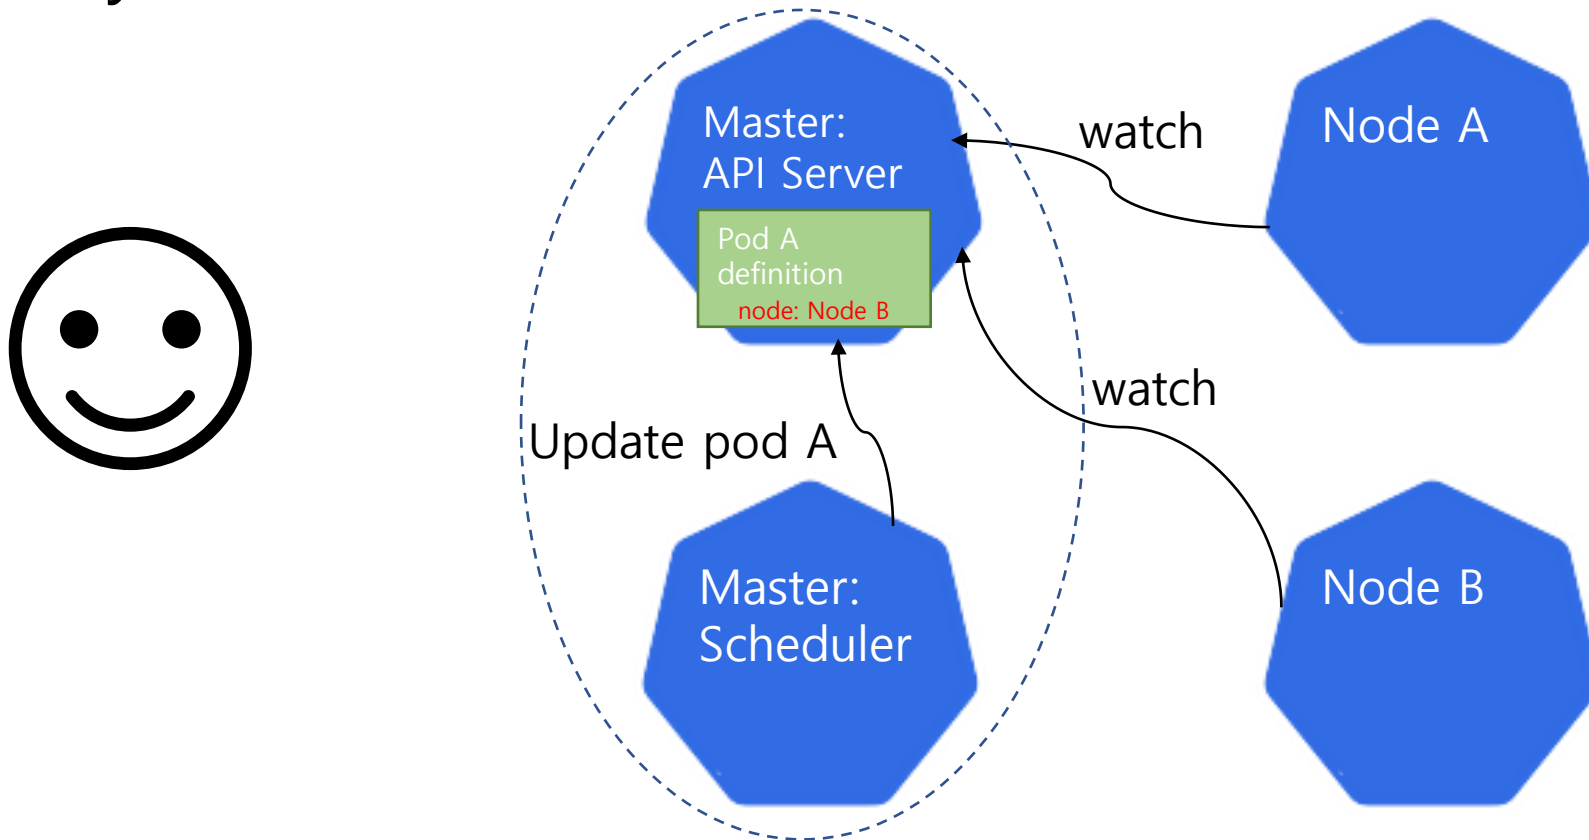
No hidden internal APIs

- All components watch the Kubernetes API, and figure out what they need to do.



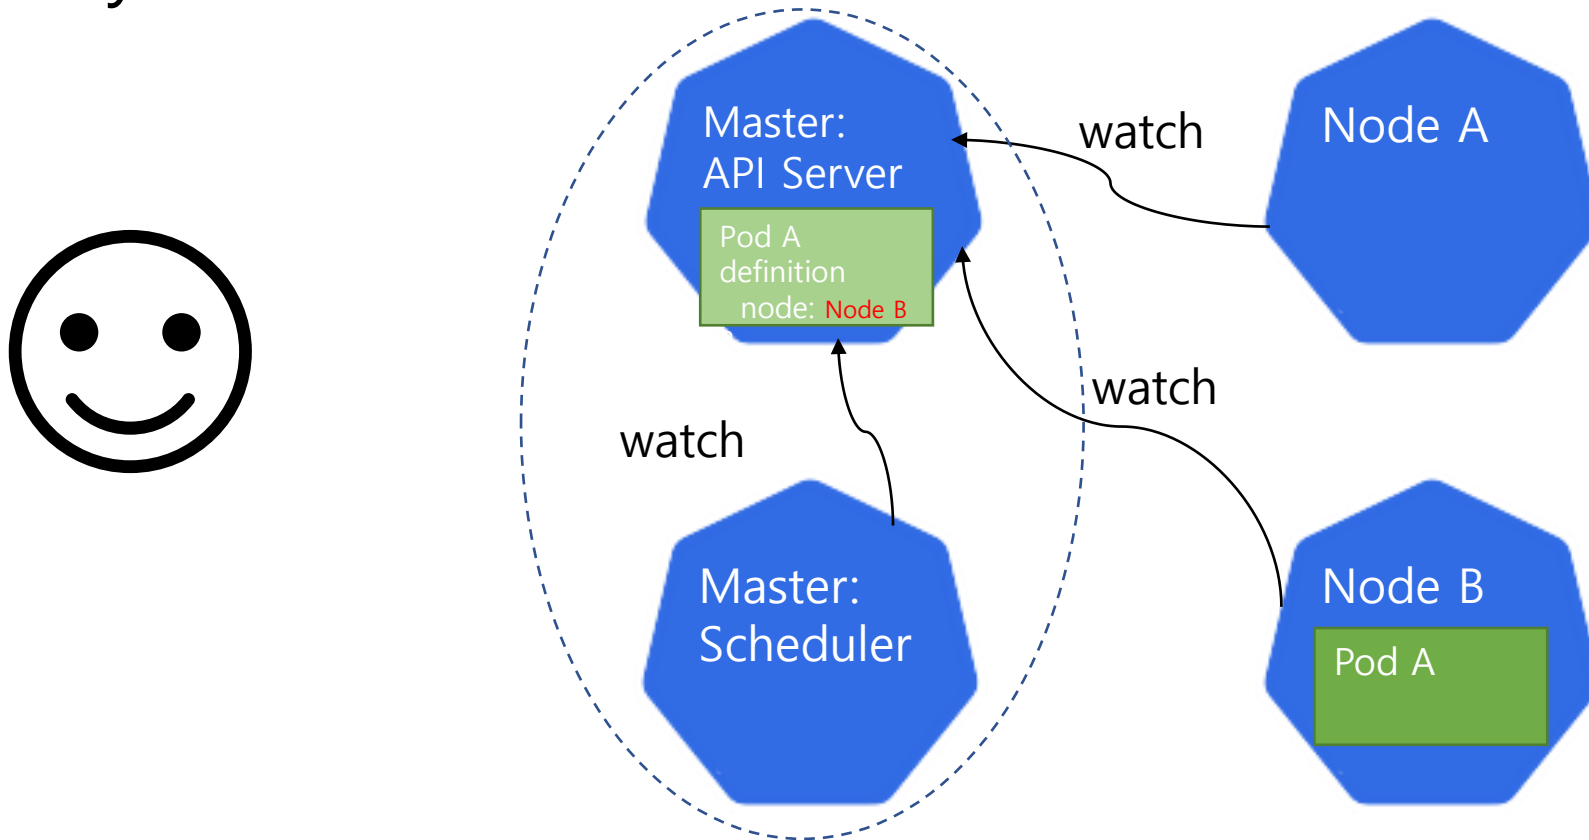
No hidden internal APIs

- All components watch the Kubernetes API, and figure out what they need to do.



No hidden internal APIs

- All components watch the Kubernetes API, and figure out what they need to do.



Why No hidden internal APIs?

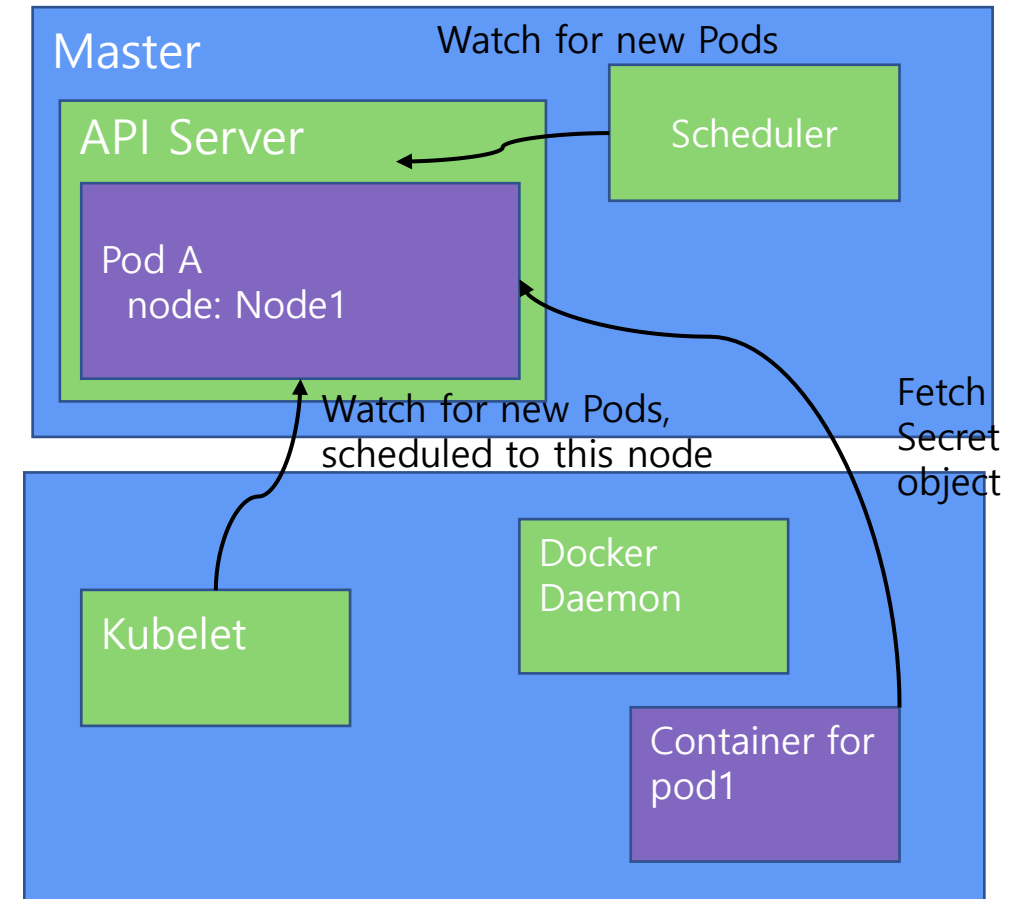
- **Declarative** API provides the same benefits to **internal components**
- Resulting in a Simpler, **more robust system** that can easily recover from failure of components.
 - No single point of failure.
 - Simple master components.
- Also **makes Kubernetes composable and extensible.**
 - Default component not working for you?
 - Turn it off and replace it with your own.
 - Additional functionality not yet available?
 - Write your own and to add it.

Kube API Data

- Kubernetes API has lots of **data** that is interesting to workloads
 - Secrets – Sensitive info stored in KubeAPI
 - e.g. passwords, certificates, etc.
 - ConfigMap – Configuration info stored in KubeAPI
 - e.g. application startup parameters, etc.
 - DownwardAPI – Pod information in KubeAPI
 - e.g. name/namespace/uid of my current pod.

Fetching Kube API data

- How does application fetch secrets, config map, etc. information?
- **Principle:** No hidden internal APIs.
- **Obvious solution:** Modify app to read directly from API Server.
 - > Not good for legacy applications
 - So, principle #3



Principle #3

**Meet the user where they are.
(i.e. Supporting Legacy Applications)**

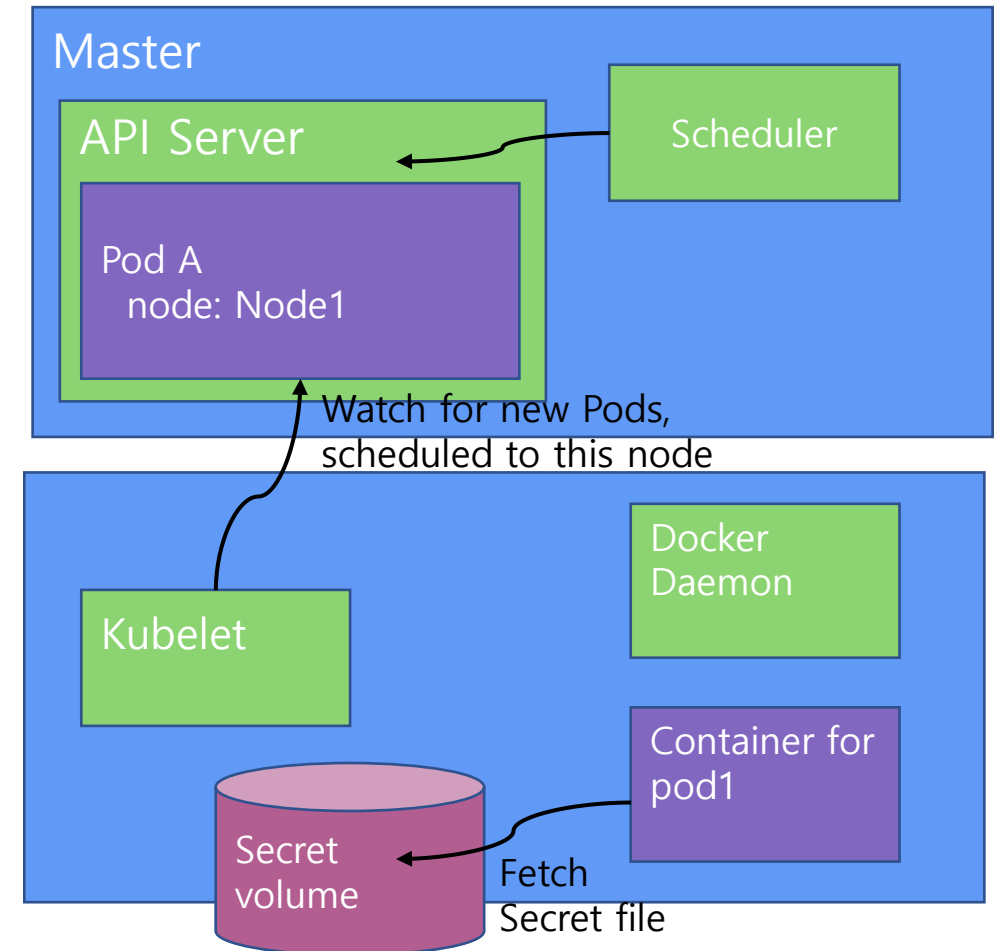
기존 응용들을 수정없이 돌릴 수 있게 하자

Why meet the user where they are?

- Minimize hurdles for deploying workloads on Kubernetes.
- Increases adoption.

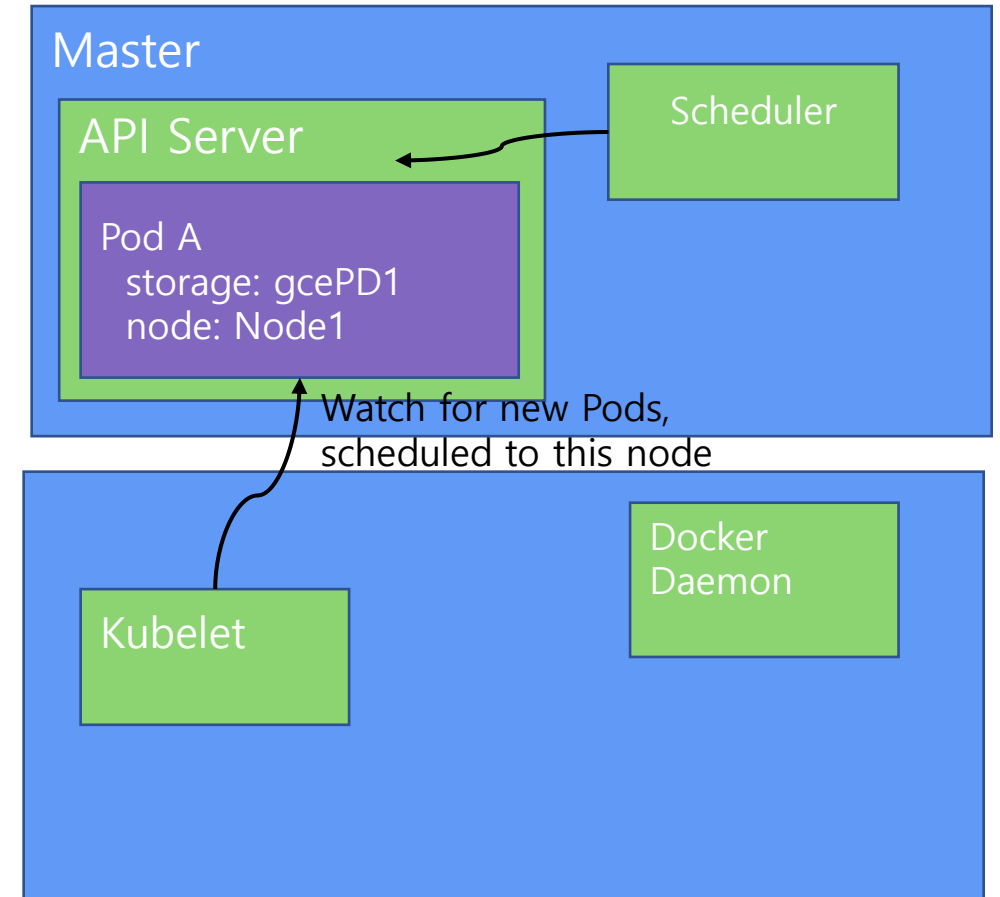
Meet the user where they are.

- Before:
 - App must be modified to be Kubernetes aware.
- After:
 - If app can load config or secret data from file or environment variables it doesn't need to be modified!
 - But, not good solution, because PoD's could be deleted after application's mission completion
 - So,... Remote storage for stable data storage



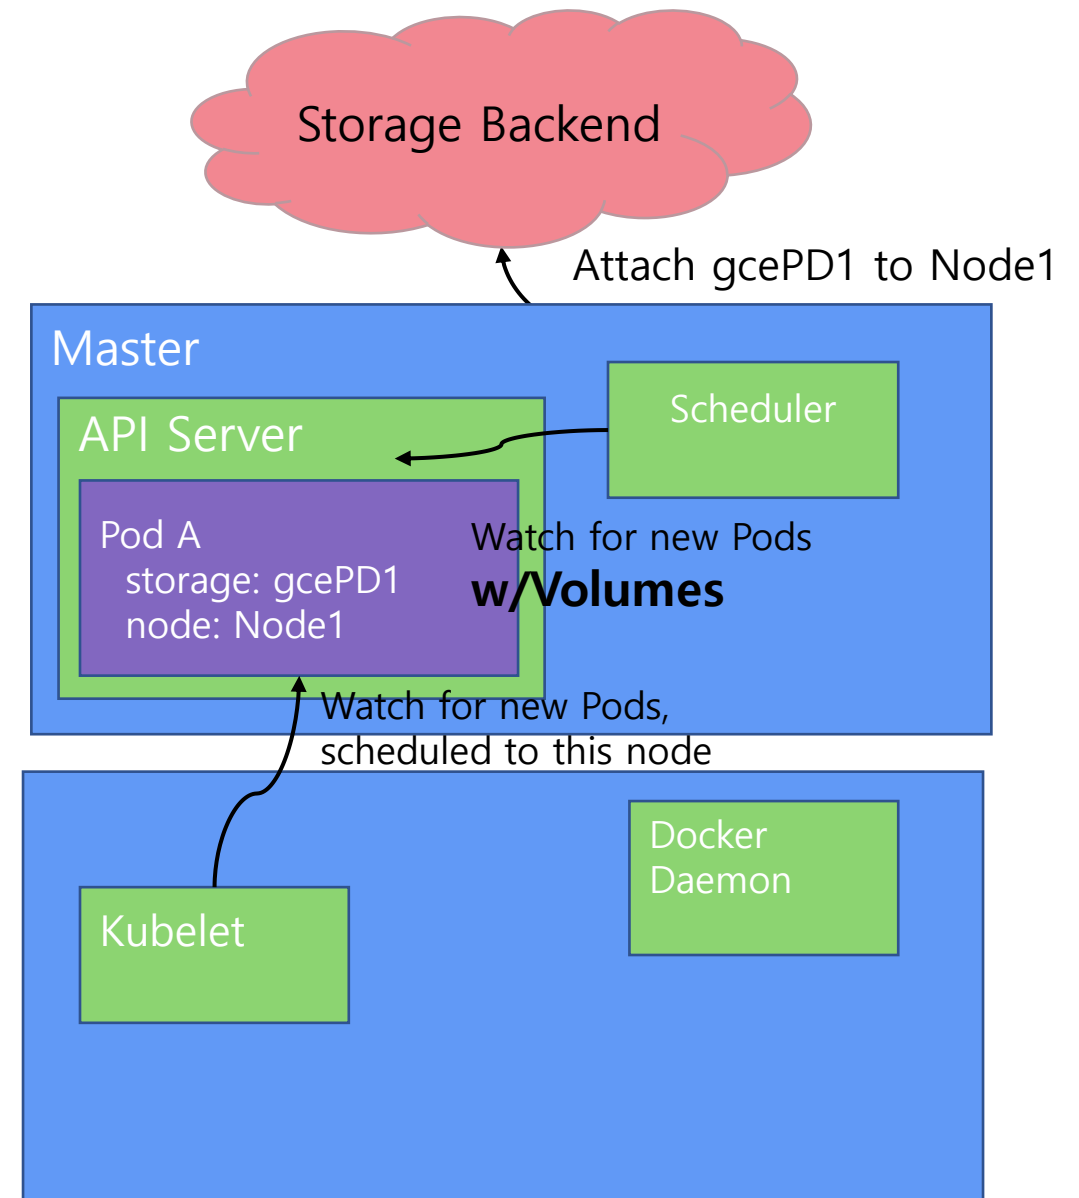
Remote Storage

- Could directly reference a remote volume (GCE PD, AWS EBS, NFS, etc.) in pod definition.
- Kubernetes will automatically make it available to workload



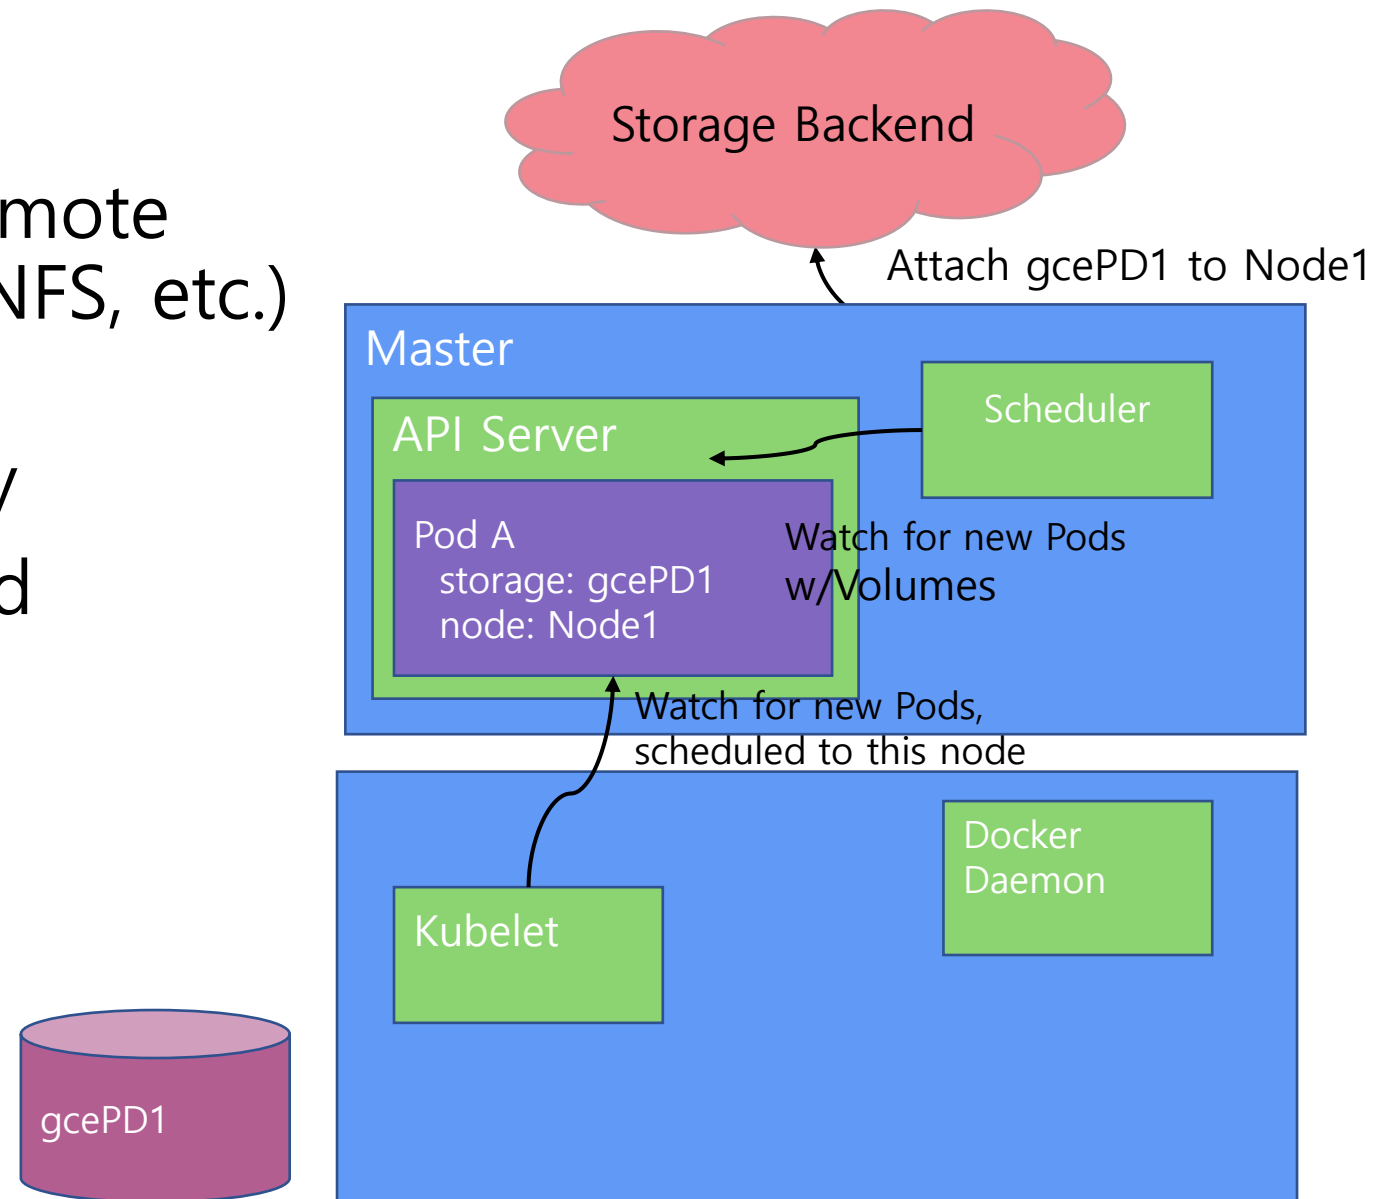
Remote Storage

- Could directly reference a remote volume (GCE PD, AWS EBS, NFS, etc.) in pod definition.
- Kubernetes will automatically make it available to workload



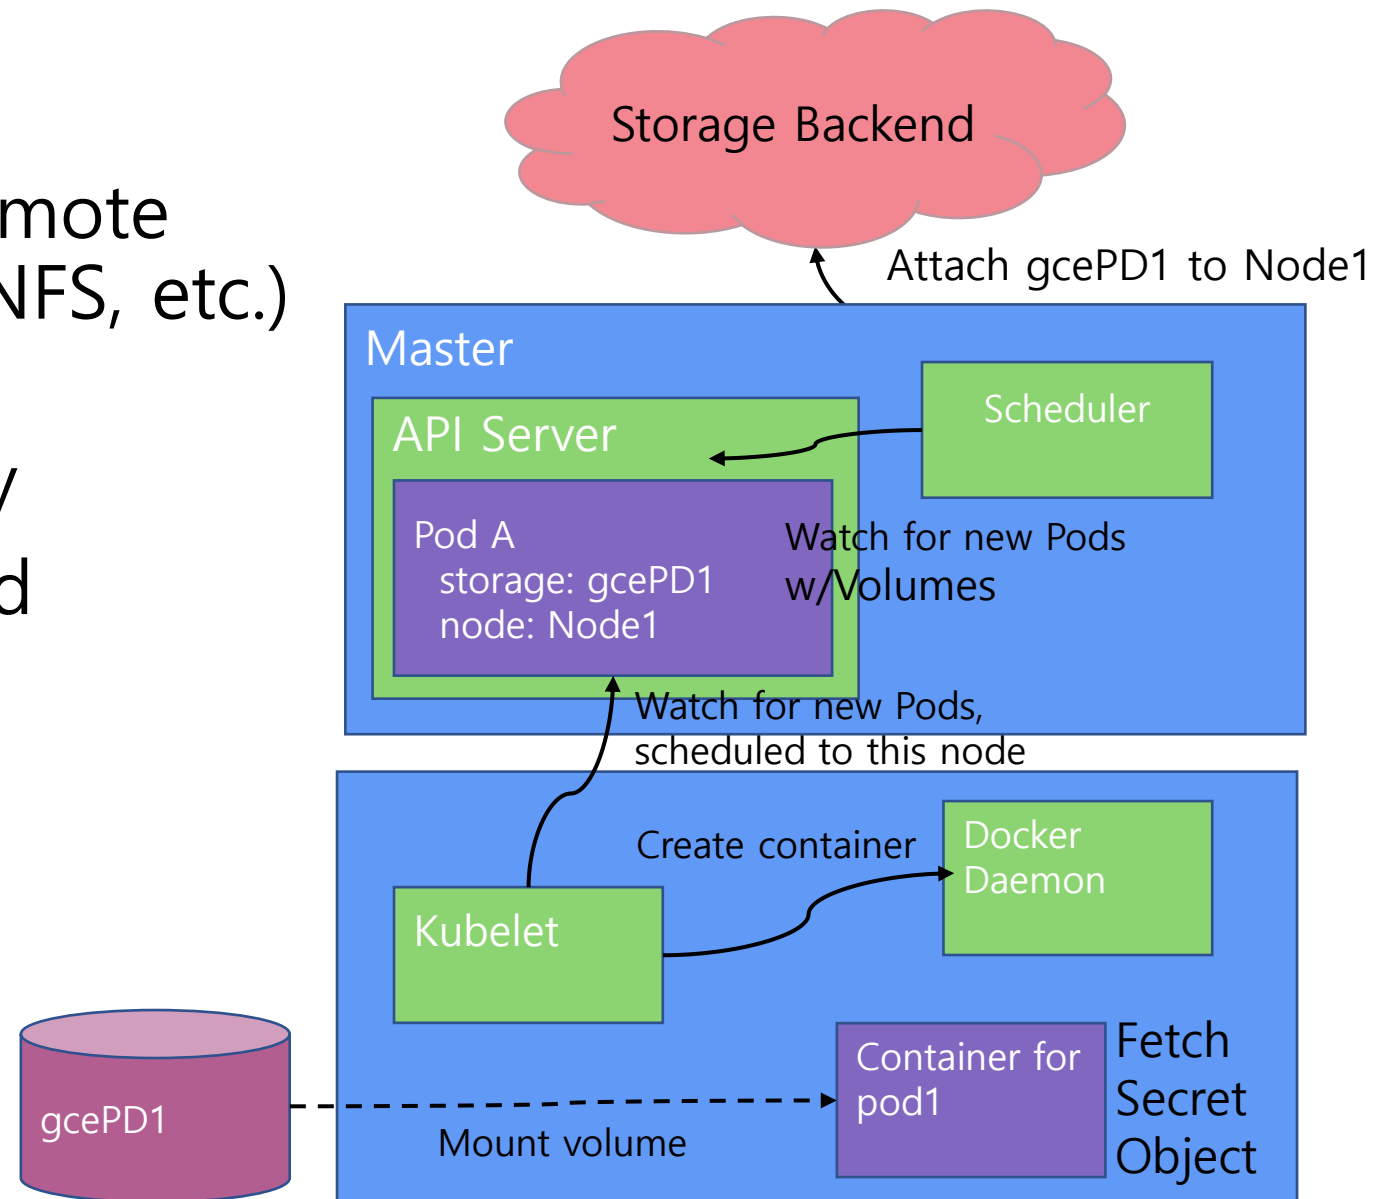
Remote Storage

- Could directly reference a remote volume (GCE PD, AWS EBS, NFS, etc.) in pod definition.
- Kubernetes will automatically make it available to workload



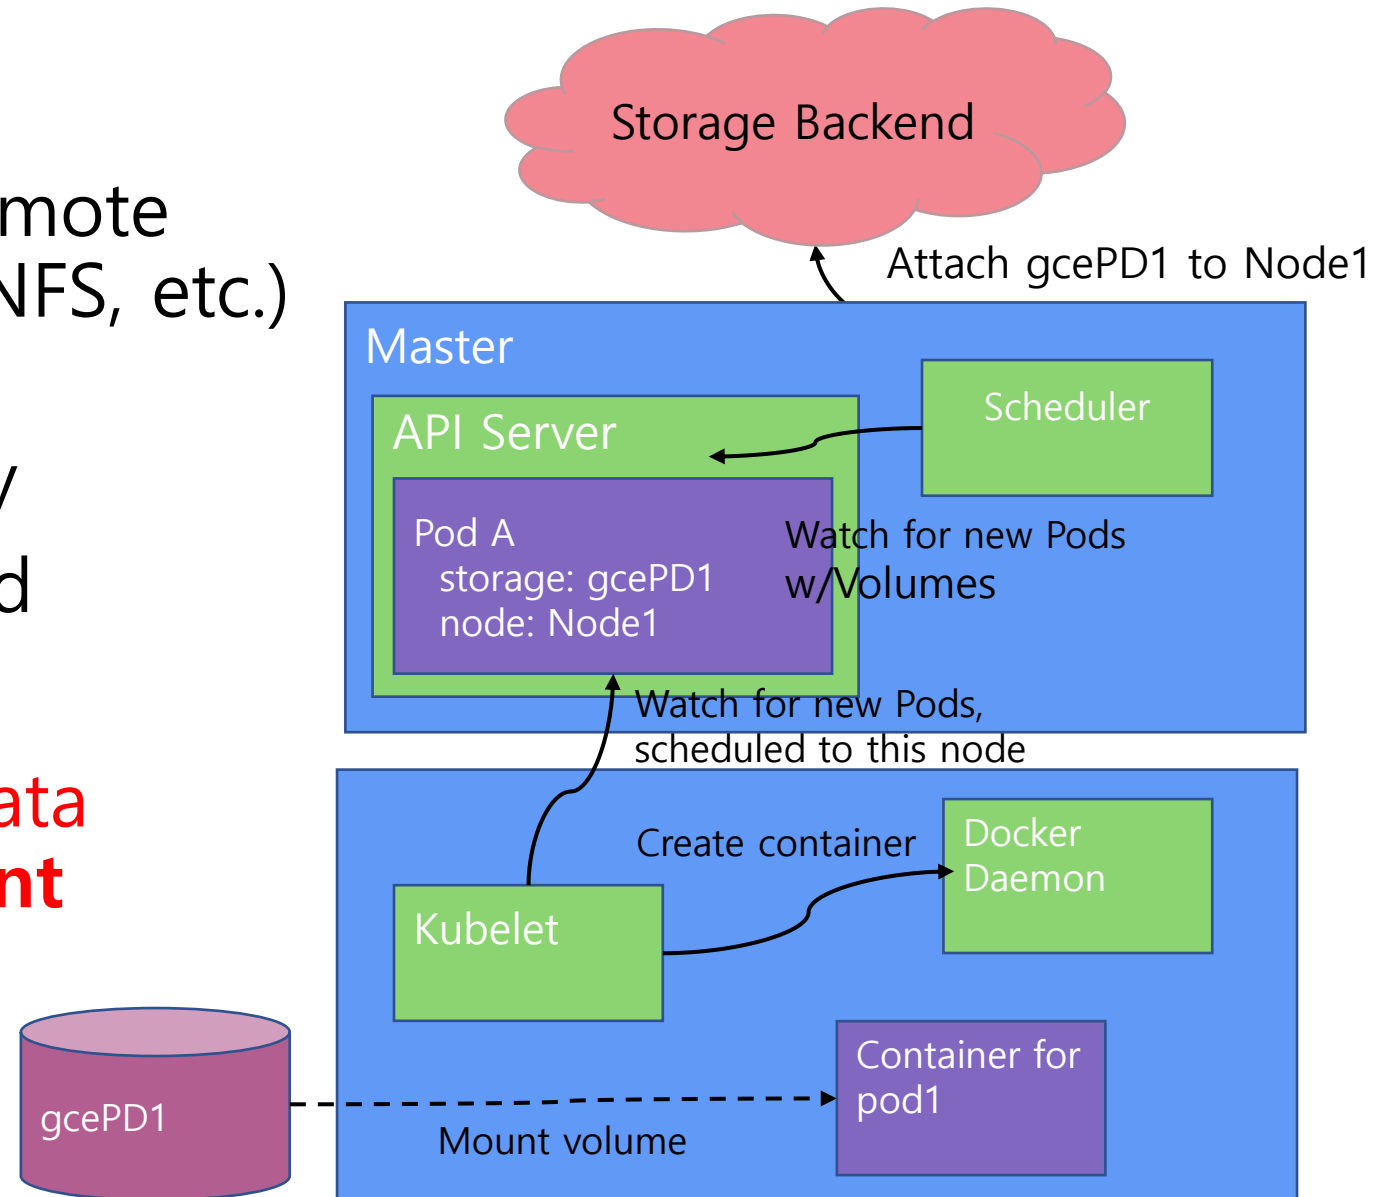
Remote Storage

- Could directly reference a remote volume (GCE PD, AWS EBS, NFS, etc.) in pod definition.
- Kubernetes will automatically make it available to workload



Remote Storage

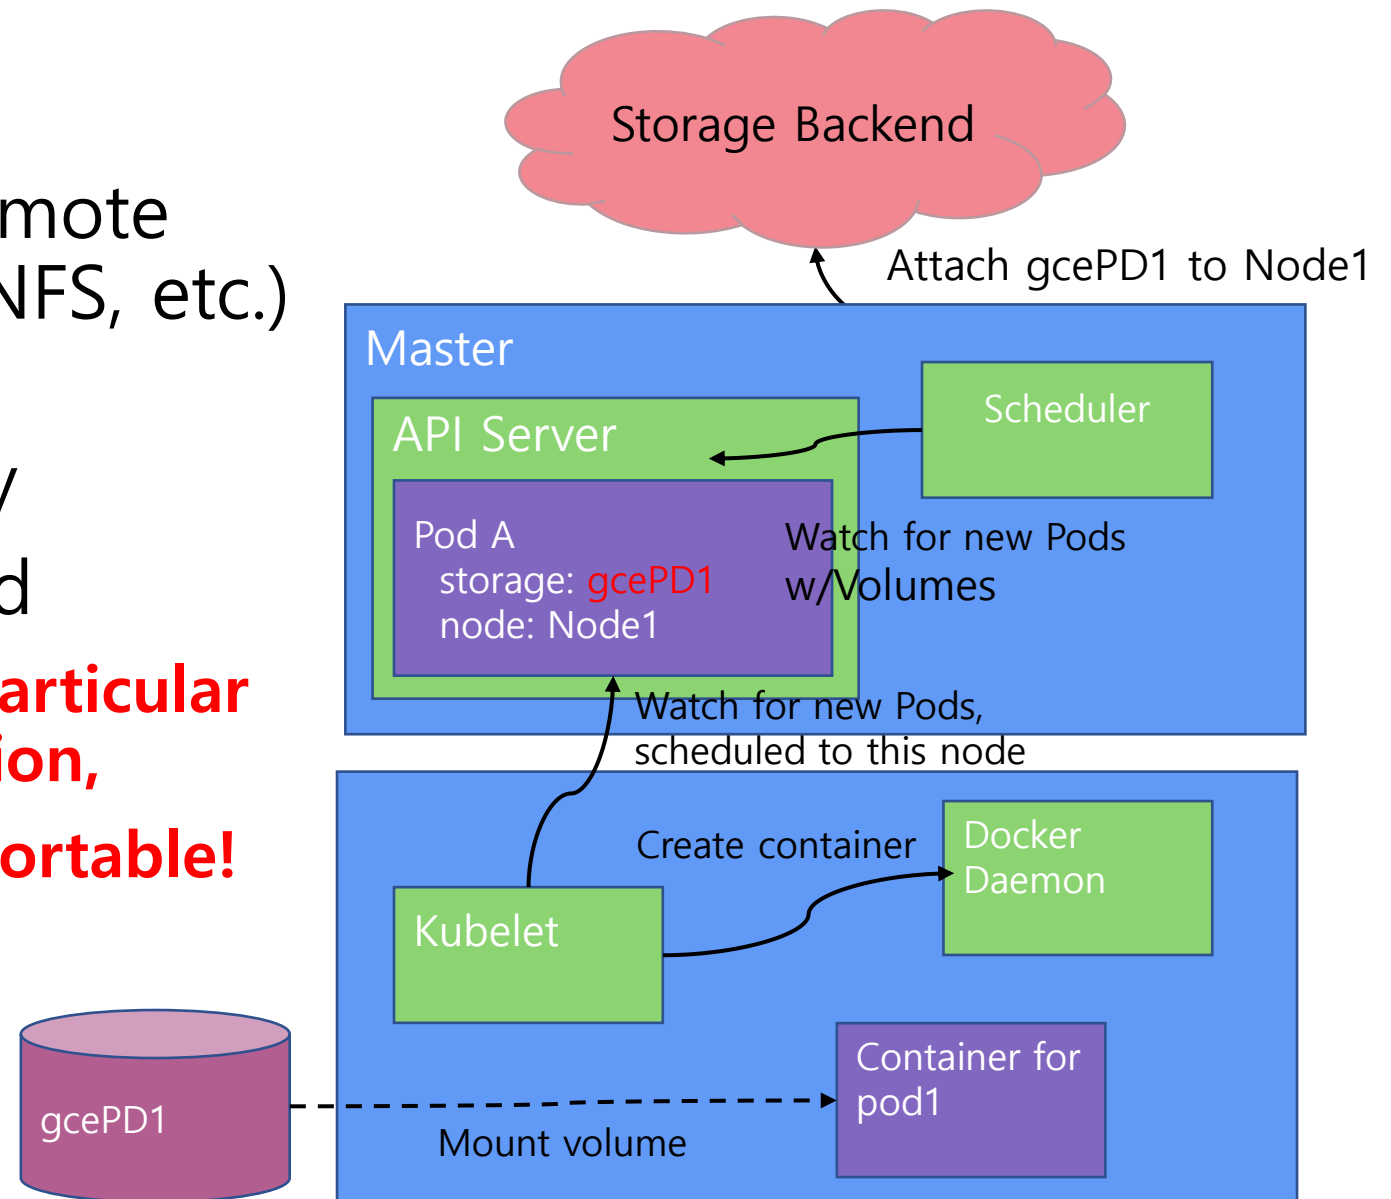
- Could directly reference a remote volume (GCE PD, AWS EBS, NFS, etc.) in pod definition.
- Kubernetes will automatically make it available to workload
- Even If PoD's deleted ..the **data** stored in Volume is **persistent**



Remote Storage

- Could directly reference a remote volume (GCE PD, AWS EBS, NFS, etc.) in pod definition.
- Kubernetes will automatically make it available to workload

But...If you directly reference a particular type of disk inline of PoD definition, the PoD definition is no longer portable!



Principle #4

Workload portability

PoD의 이식성을 강화(스토리지 측면에서도)

PVC/PV

- PersistentVolume and PersistentVolumeClaim Abstraction
- **Decouple storage implementation from storage consumption**

Why Workload Portability?

- Decouple distributed system **application development** from **cluster implementation**
- **Make Kubernetes a true abstraction layer, like an OS.**

Kubernetes Principles introduced

- 1. Kube API declarative over imperative.
- 2. No hidden internal APIs
- 3. Meet the user where they are: Remote storage
- 4. Workload portability: PV/PVC