 Soongsil Univ.
<http://design.ssu.ac.kr>

 Digital
System
Design Lab.

Ch. 03 Processes


Chanho Lee
Soongsil University

Copyright 2023. () all rights reserved

1

Objectives

- Components of a process : representation and scheduling
- Process creation and termination and using the system calls
- Interprocess communication(IPC) using shared memory and message passing.
 - ▶ Pipes and POSIX shared memory
 - ▶ client-server communication using sockets and remote procedure calls (RPC)
- Design kernel modules in Linux

 Soongsil Univ.
Digital System Design Lab.

Operating Systems

Copyright 2023. () all rights reserved 2

2


Soongsil Univ.
http://design.ssu.ac.kr


Digital
System
Design Lab.

4주차 process concept

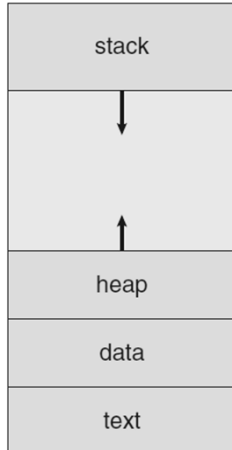
Copyright 2023. () all rights reserved

3


3.1 Process Concept

- Processes
 - ▶ Jobs (batch system) tasks or user programs (time-shared system)
 - ▶ Internal activities of OS
 - ▶ Program(executable file, passive) in execution: active
 - ▶ The status of the current activity of a process
 - program counter + processor's registers
 - ▶ memory layout of a process
 - Text—the executable code
 - Data—global variables
 - Heap—dynamically allocated during run time
 - Stack—temporary data storage (such as function parameters, return addresses, and local variables activation record)
 - ▶ an execution environment for other code
 - Ex. JAVA: an executable Java program within the Java virtual machine (JVM)

max



0


Soongsil Univ.
Digital System Design Lab.

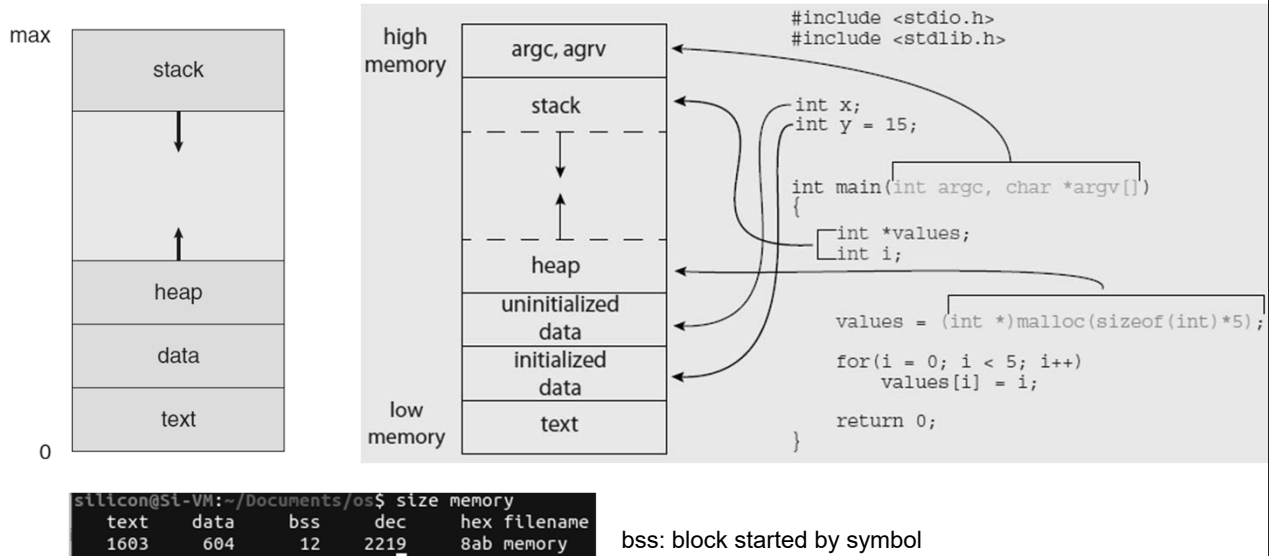
Operating Systems

Copyright 2023. () all rights reserved 4

4

3.1 Process Concept

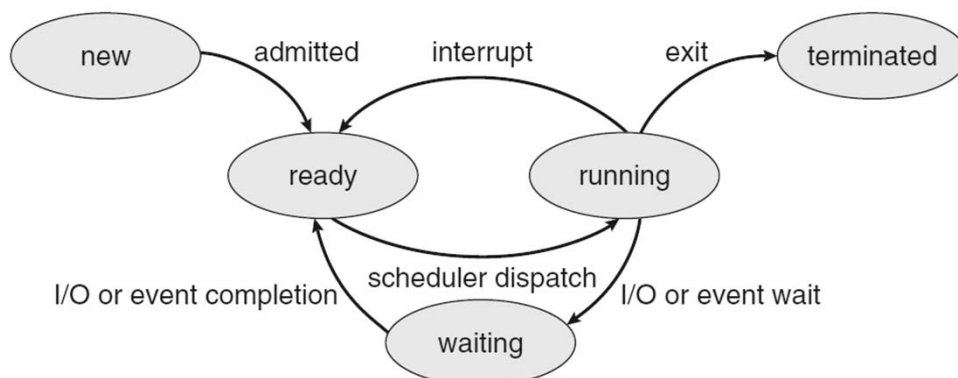
● Memory layout of a C program



3.1 Process Concept

● Process State

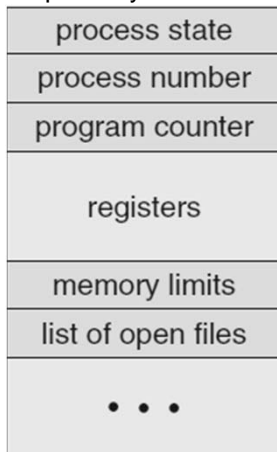
► defined in part by the current activity of that process



3.1 Process Concept

● Process Control Block (PCB)

- ▶ task control block (TCB)
- ▶ contains many pieces of information associated with a specific process
- ▶ repository for all the data needed to start, or restart, a process, along with accounting data



address of the next instruction

CPU-scheduling information: process priority, pointers to scheduling queues, and any other scheduling parameters. (ch.5)

Memory-management information: value of the base and limit registers and the page tables, or the segment tables, and so on. (ch.9).

Accounting information: the amount of CPU and real time used, time limits, account numbers, job or process numbers, and so on.

I/O status information: list of I/O devices allocated to the process, a list of open files, and so on.

3.1 Process Concept

● Threads (ch.4)

- ▶ Single thread of execution: a process to perform only one task at a time
 - ex. Word-processor: typing and then spell checking
- ▶ Multiple threads of execution: a process to perform multiple tasks at a time
 - beneficial on multicore system
 - ex. Word-processor: typing and spell checking simultaneously (2 threads)
- ▶ PCB includes thread information for multiple threads

3.2 Process Scheduling

● Process scheduler

- ▶ selects an available process

The objective of multiprogramming: to maximize CPU utilization.
The objective of time sharing: frequent switching (multitasking)

- ▶ Multiprogramming: # of process > # of cores

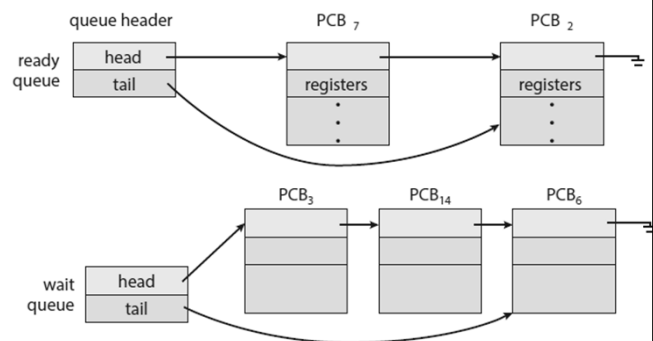
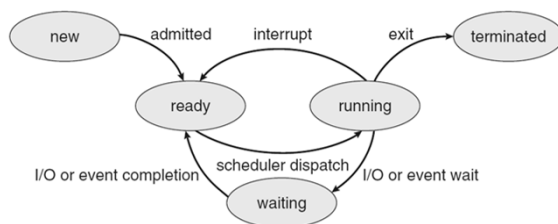
- Each CPU core can run one process at a time: waiting (ready) processes
- Degree of multiprogramming: the number of processes currently in memory

- ▶ CPU-bound or I/O bound process

● Scheduling Queues

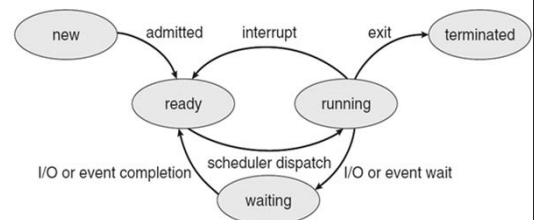
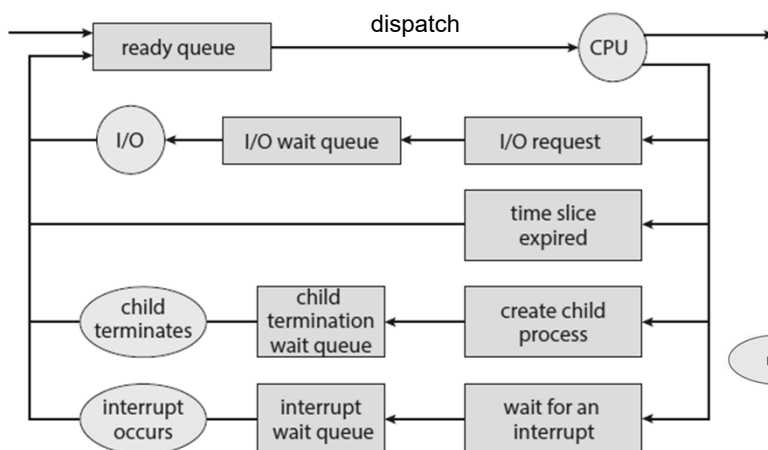
- ▶ Ready queue: processes ready to run

- ▶ Wait queue: processes waiting for events



3.2 Process Scheduling

● Queueing diagram



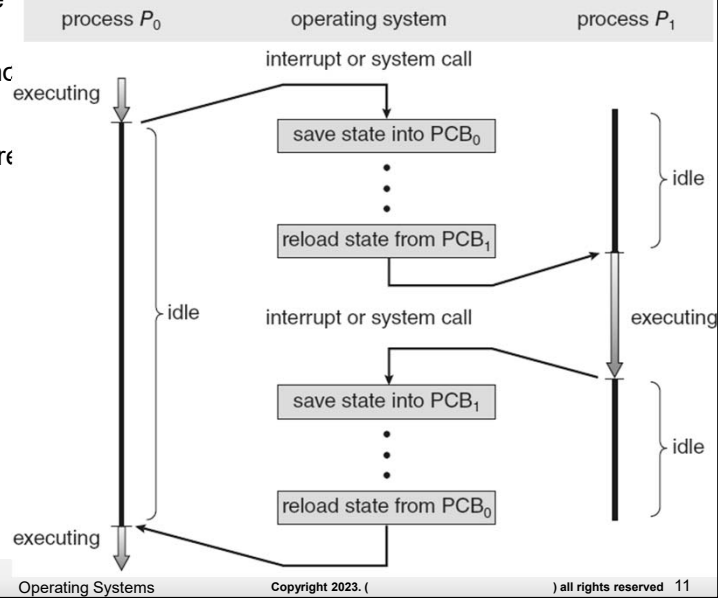
3.2 Process Scheduling

● CPU Scheduling

- ▶ CPU scheduler: selects a process in the
 - I/o-bound and CPU-bound process
- ▶ Swapping: remove a process from memc

● Context Switch

- ▶ Context: includes the value of the CPU r
information in PCB
- ▶ State save and restore ✓
- ▶ Context-switch time is pure overhead
 - Typically, a several μ s



11

3.3 Operations on Processes

● Process Creation

- ▶ A process may create several new pro
- ▶ Unique process identifier (or pid): inde
- ▶ Parent-child
- ▶ Child: copy of PCB from parent

the first
when

```
silicon@silicon-VirtualBox:~$ pstree
systemd--ModemManager--2*[{ModemManager}]
systemd--NetworkManager--2*[{NetworkManager}]
systemd--accounts-daemon--2*[{accounts-daemon}]
systemd--acpid
systemd--anacron
systemd--avahi-daemon--avahi-daemon
systemd--canonical-livep--12*[{canonical-livep}]
systemd--colord--2*[{colord}]
systemd--cron
systemd--cups-browsed--2*[{cups-browsed}]
systemd--cupsd
systemd--dbus-daemon
systemd--fwupd--4*[{fwupd}]
systemd--gdm3--gdm-session-wor--gdm-x-session--Xorg--{Xorg}
systemd--gdm3--gdm-session-wor--gdm-x-session--gnome-session-b--ssh-agent--2*[{gnome-}
systemd--gdm3--gdm-session-wor--gdm-x-session--2*[{gdm-x-session}]
systemd--gnome-keyring-d--3*[{gnome-keyring-d}]
systemd--irqbalance--{irqbalance}
systemd--2*[kerneloops]
systemd--networkd-dispat
```



Figure 3.7 A tree of processes on a typical Linux system.

12

3.3 Operations on Processes

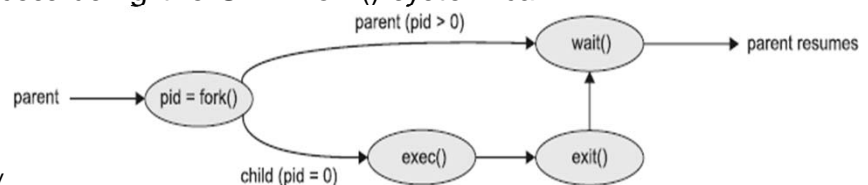
● Process Creation

- ▶ Resources of child
 - From OS
 - From parent: constrained to a subset
- ▶ Execution of parent
 - Concurrent with child
 - Waiting until the termination of child
- ▶ Address-space possibilities for child
 - duplicate of the parent
 - new program loaded

3.3 Operations on Processes

● Creating a separate process using the UNIX fork() system call

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
int main()
{
    pid_t pid;
    /* fork a child process */
    pid = fork(); //pid(p)>0: pid of child process; pid(c): 0
    if (pid < 0) { /* error occurred. Child starts here */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL); // child ≠ parent now
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL); // transit to wait state
        printf("Child Complete");
    }
    return 0;
}
```



3.3 Operations on Processes

● Process Termination

▶ exit() system call in child

- resources—including physical and virtual memory, open files, and I/O buffers—are deallocated
- return a status value to waiting parent process (via wait())
 - Ex: exit status = 1

```
/* exit with status 1 */
exit(1); //explicitly or implicitly
```

```
– Ex. pid_t pid;
      int status;
      pid = wait(&status); //parent gets exit status and child's pid
```

- Zombie: child terminated before wait() is called usually temporary
- Orphan: parent terminated before calling wait() init(systemd) is assigned as a new parent in UNIX/Linux

▶ Termination by another process

- By parent or OS
 - Excessive usage of resources by child
 - Task of child is no longer required
 - Parent is exiting or terminated abnormally (cascading termination)

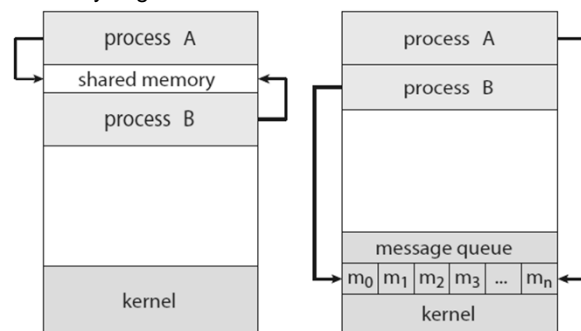
3.4 Interprocess Communication

● Sharing data or not: independent or cooperating

- Information sharing: concurrent access to information.
- Computation speedup: breaking into subtasks, executing in parallel with multiple cores.
- Modularity: dividing system functions into separate processes or threads.

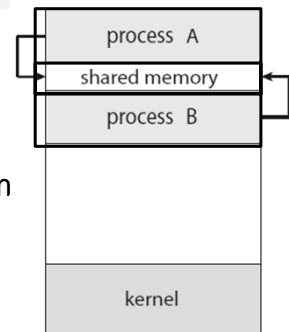
● Interprocess communication (IPC)

- shared memory: faster than message passing
 - system calls are required only to establish shared memory regions
- message passing
 - Typically implemented using system calls
 - useful for exchanging smaller amounts of data
 - easier to implement in a distributed system



3.5 IPC in Shared-Memory Systems

- Shared-memory
 - ▶ requires communicating processes to establish
 - ▶ Remove protection for shared memory
 - ▶ Read and write without OS's control
- Example of cooperating processes: producer-consumer problem
 - ▶ Concurrent operation of producer and consumer
 - ▶ a buffer of items in shared memory
 - ▶ filled by the producer and emptied by the consumer: synchronized
 - ▶ unbounded buffer/bounded buffer



3.5 IPC in Shared-Memory Systems

- Producer-consumer with bounded buffer

```
#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;
item buffer[BUFFER_SIZE]; //circular array with two logical pointers: in, out
int in = 0; // empty when in=out
int out = 0; // full when ((in + 1) % BUFFER_SIZE) == out
```

```
item next_produced;
while (true) {
    /* produce an item in next_produced */
    while (((in + 1) % BUFFER_SIZE) == out) //full
        ; /* do nothing */
    buffer[in] = next_produced; //write
    in = (in + 1) % BUFFER_SIZE;
} //max BUFFER_SIZE-1 items in BUFFER: why?
```

```
item next_consumed;
while (true) {
    while (in == out) //empty
        ; /* do nothing */
    next_consumed = buffer[out]; //read
    out = (out + 1) % BUFFER_SIZE;
    /* consume the item in next_consumed */
}
```

3.6 IPC in Message-Passing Systems

● Message passing

- ▶ to communicate and to synchronize their actions without sharing the same address space.
- ▶ Operations: `send(message)` and `receive(message)`
- ▶ fixed or variable in size
- ▶ communication link:
 - implemented physically in a variety of ways
 - Logical implementation.
 - Direct or indirect communication
 - Synchronous or asynchronous communication
 - Automatic or explicit buffering

3.6 IPC in Message-Passing Systems

● Naming

- ▶ Direct (symmetry)
 - `send(P, message)`—Send a message to process P.
 - `receive(Q, message)`—Receive a message from process Q.
 - A link is established automatically using identity
 - A link is associated with exactly two processes
 - Between each pair of processes, there exists exactly one link.
- asymmetry
 - `send(P, message)`—Send a message to process P.
 - `receive(id, message)`—Receive a message from any process
- limited modularity: changing the identifier requires examining all other definitions

3.6 IPC in Message-Passing Systems

► Indirect

- messages are sent to and received from mailboxes, or ports
- Mailbox: an object with a unique identification containing messages: placed and removed.
- send(A, message)—Send a message to mailbox A. (A: integer for POSIX)
- receive(A, message)—Receive a message from mailbox A.
- A link is established with a shared mailbox
- A link may be associated with more than two processes
- Between each pair of processes, a number of different links with corresponding mailboxes are allowed
- Mailboxes may be owned by a process or OS

3.6 IPC in Message-Passing Systems

● Synchronization

- Blocking send: sending is blocked until the message is received or by the mailbox.
- Nonblocking send: sends the message and resumes operation.
- Blocking receive: receiver blocks until a message is available.
- Nonblocking receive: receiver retrieves either a valid message or a null.
- combinations of send() and receive() are possible
 - Both send() and receive() are blocking, we have a rendezvous: trivial solution to the producer-consumer problem

```
message next produced;
while (true) {
    /* produce an item in next produced */
    send(next produced);
}
```

```
message next consumed;
while (true) {
    receive(next consumed);
    /* consume the item in next consumed */
}
```

3.6 IPC in Message-Passing Systems

● Buffering: temporary queue implementation for IPC

- ▶ Zero capacity:
 - maximum length of zero:
 - no waiting messages.
 - blocking only.
- ▶ Bounded capacity.
 - finite length n: max n messages.
 - the sender is blocked if the link is full
- ▶ Unbounded capacity:
 - potentially infinite length:
 - the sender is never blocked

3.7 Examples of IPC Systems

● POSIX Shared Memory

▶ using memory-mapped files

- `fd = shm_open(name, O_CREAT | O_RDWR, 0666); //Create a shared-memory object`

Integer file descriptor

Created if it does not yet exist

for reading and writing

file-access permissions

- `ftruncate(fd, 4096); //to configure the size of the object in bytes`
- `mmap()` function establishes a memory-mapped file containing the shared-memory object

3.7 Examples of IPC Systems

```
int main() //producer
{
    const int SIZE = 4096; //the size (in bytes) of shared memory object
    const char *name = "OS"; // name of the shared memory object
    /* strings written to shared memory */
    const char *message_0 = "Hello";
    const char *message_1 = "World!";
    int fd; //shared memory file descriptor
    char *ptr; //pointer to shared memory object
    fd = shm_open(name,O_CREAT | O_RDWR,0666); //create the shared memory object
    ftruncate(fd, SIZE); //configure the size of the shared memory object
    /* memory map the shared memory object */
    ptr = (char *)mmap(0, SIZE, PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0); //read and write
    /* write to the shared memory object */
    sprintf(ptr,"%s",message_0); //write formatted string to the shared-memory object
    ptr += strlen(message_0);
    sprintf(ptr,"%s",message_1);
    ptr += strlen(message_1);
    return 0;
}
```

start

offset

changes to the shared-memory object will be visible

gcc -o producer producer.c -lrt

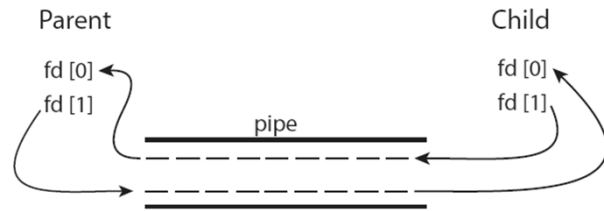
3.7 Examples of IPC Systems

```
int main() //consumer
{
    const int SIZE = 4096; //the size (in bytes) of shared memory object
    const char *name = "OS"; // name of the shared memory object
    /* strings written to shared memory */
    int fd; //shared memory file descriptor
    char *ptr; //pointer to shared memory object
    fd = shm_open(name, O_RDONLY, 0666); //open the shared memory object
    /* memory map the shared memory object */
    ptr = (char *)mmap(0, SIZE, PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0);
    /* read from the shared memory object */
    printf("%s",(char *)ptr);
    shm_unlink(name); //remove the shared memory object
    return 0;
}
```

3.7 Examples of IPC Systems

● Pipes

- ▶ simple ways for IPC
- ▶ Implementation issues
 - Bidirectional or unidirectional
 - half duplex or full duplex
 - Relationship (such as parent-child) required?
 - communicate over a network or not?
- ▶ Ordinary Pipes
 - Unidirectional
 - cannot be accessed from outside; typically, IPC between parent-child
 - child inherits the pipe
 - Two pipes for two-way communication
 - pipe(int fd[]): fd[0]: read end, and fd[1]: write end
 - special type of file: read() and write() system calls
 - exist only while the processes are communicating with one another



3.7 Examples of IPC Systems

```
#define BUFFER_SIZE 25
#define READ_END 0
#define WRITE_END 1
int main(void)
{
    char write_msg[BUFFER_SIZE] = "Greetings";
    char read_msg[BUFFER_SIZE];
    int fd[2];
    pid_t pid;
    if (pipe(fd) == -1) { // create the pipe
        fprintf(stderr, "Pipe failed");
        return 1;
    }
    pid = fork(); // fork a child process
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
}
```

```
if (pid > 0) { /* parent process */
    close(fd[READ_END]); // close the unused end of the pipe
    /* write to the pipe */
    write(fd[WRITE_END], write_msg, strlen(write_msg)+1);
    /* close the write end of the pipe */
    close(fd[WRITE_END]);
}
else { /* child process */
    /* close the unused end of the pipe */
    close(fd[WRITE_END]);
    /* read from the pipe */
    read(fd[READ_END], read_msg, BUFFER_SIZE);
    printf("read %s", read_msg);
    /* close the read end of the pipe */
    close(fd[READ_END]);
}
return 0;
}
```

3.7 Examples of IPC Systems

► Named Pipes

- Bidirectional (half-duplex), and no parent–child relationship is required
- several processes can use it for communication
- continue to exist after communicating processes have finished
- referred to as FIFOs in UNIX systems
- UNIX/Linux
 - `mkfifo()` system call
 - manipulated with the ordinary `open()`, `read()`, `write()`, and `close()` system calls
 - byte-oriented data may be transmitted

PIPES IN PRACTICE

`ls | less`

: results of list directory → input of less (listing a screen at a time)

3.8 Communication in Client–Server Systems

● Shared memory and message passing +

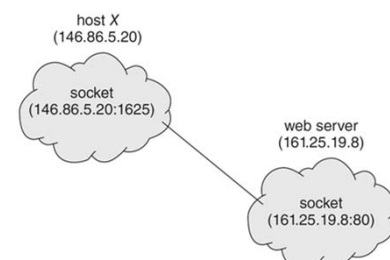
● Sockets

► Socket:

- endpoint for communication
- identified by an IP address concatenated with a port number : 146.86.5.20:1625
- client–server architecture

► A pair of processes communicating over a network employs a pair of sockets

- Server waits for incoming client requests by listening to a specified port.
- Servers implementing specific services listen to well-known ports (<1024)
 - SSH:22; FTP: 21; HTTP: 80
- client process is assigned a port by its host: arbitrary number > 1024
- All connections must be unique: different port



3.8 Communication in Client–Server Systems

- Remote Procedure Calls (RPC)

- ▶ One of the most common forms of remote service
- ▶ well structured messages: no longer just packets of data
- ▶ contains an identifier specifying the function and parameters: execution and output sent back
- ▶ One IP address with multiple ports : multiple services
- ▶ Useful in implementing a distributed file system

Summary

- Process: a program in execution, the status of the current activity
- Layout of a process in memory: (1) text, (2) data, (3) heap, and (4) stack
- Process states: (1) ready, (2) running, (3) waiting, and (4) terminated
- PCB: kernel data structure
- Process scheduler: to select a process to run on a CPU.
- Context switch: switches from running one process to running another
- The fork(): to create processes
- IPC: shared memory/exchanging messages/pipe
- Client–server communication
 - ▶ Sockets: data packet
 - ▶ RPC: abstracting the concept of function (procedure) calls on a separate computer

Exercises, problems and projects

● Exercises

- ▶ 3.1, 3.2, 3.5,
- ▶ 3.10, 3.11, 3.12, 3.13, 3.16

● Problems

▶ 3.4

3.10 Explain the role of the `init` (or `systemd`) process on UNIX and Linux systems in regard to process termination.

● Projects

3.11 Including the initial parent process, how many processes are created by the program shown below.

```
#include <stdio.h>
#include <unistd.h>
int main()
{
    int i;
    for (i = 0; i < 4; i++)
        fork();
    return 0;
}
```

3.12 Explain the circumstances under which the line of code marked `printf("LINE J")` will be reached.

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
int main()
{
    pid_t pid;
    pid = fork(); /* fork a child process */
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
        printf("LINE J");
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }
    return 0;
}
```

Exercises, problems and projects

3.13 Using the program in Figure 3.34, identify the values of `pid` at lines A, B, C, and D. (Assume that the actual pids of the parent and child are 2600 and 2603, respectively.)

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
int main()
{
    pid_t pid, pid1; /* fork a child process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        pid1 = getpid();
        printf("child: pid = %d", pid); /* A */
        printf("child: pid1 = %d", pid1); /* B */
    }
    else { /* parent process */
        pid1 = getpid();
        printf("parent: pid = %d", pid); /* C */
        printf("parent: pid1 = %d", pid1); /* D */
        wait(NULL);
    }
    return 0;
}
```

3.16 Using the program shown in Figure 3.35, explain what the output will be at lines X and Y.

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
#define SIZE 5
int nums[SIZE] = {0,1,2,3,4};
int main()
{
    int i;
    pid_t pid;
    pid = fork();
    if (pid == 0) {
        for (i = 0; i < SIZE; i++) {
            nums[i] *= -i;
            printf("CHILD: %d ", nums[i]); /* LINE X */
        }
    }
    else if (pid > 0) {
        wait(NULL);
        for (i = 0; i < SIZE; i++)
            printf("PARENT: %d ", nums[i]); /* LINE Y */
    }
    return 0;
}
```