
 Soongsil Univ.
<http://design.ssu.ac.kr>

 Digital
System
Design Lab.

Ch. 05 CPU Scheduling

Chanho Lee
Soongsil University

Copyright 2023. (차세대반도체 혁신공유대학 사업단) all rights reserved

1

Objectives

- Describe various CPU scheduling algorithms.
- Scheduling criteria.
- Multiprocessor and multicore scheduling.
- Real-time scheduling algorithms.
- Windows, Linux, and Solaris operating systems.
- Modeling and simulations to evaluate CPU scheduling algorithms.
- Implements several different CPU scheduling algorithms.

2


Soongsil Univ.
<http://design.ssu.ac.kr>


Digital System Design Lab.

7주차

Copyright 2023. (차세대반도체 혁신공유대학 사업단) all rights reserved

3

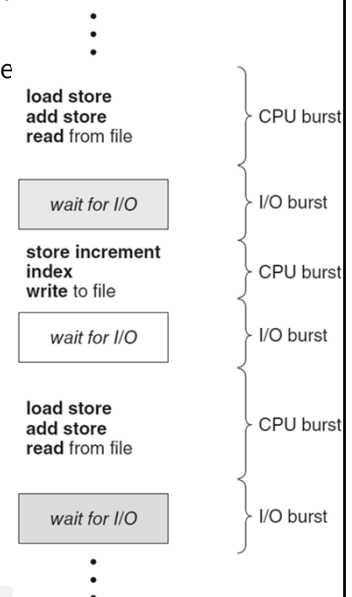
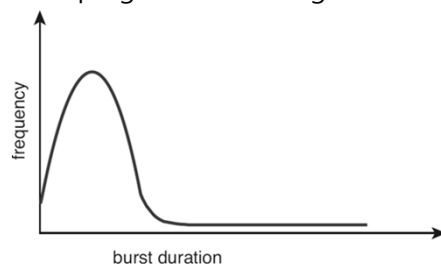
5.1 Basic Concepts

- Objective of multiprogramming: to maximize CPU utilization.

- ▶ Process execution until waiting state: typically, I/O request.
- ▶ Multiple processes in memory (ready queue) → scheduling before use
- ▶ Fundamental OS function

- CPU-I/O Burst Cycle

- ▶ Process execution: alternating sequence of CPU and I/O bursts
- ▶ Histogram of CPU-burst durations
- ▶ I/O-bound program: many short CPU bursts
- ▶ CPU-bound program: a few long CPU bursts



4

5.1 Basic Concepts

● CPU Scheduler

- ▶ selects a process in ready queue and allocates a CPU
- ▶ Ready queue: a FIFO queue, a priority queue, a tree, or simply an unordered linked list
- ▶ PCBs have the records in the queues

● Preemptive and Nonpreemptive Scheduling

- ▶ CPU-scheduling decisions when a process
 - 1. switches from the running to the waiting state or an invocation of wait()
 - 2. switches from the running to the ready state (for example, an interrupt)
 - 3. switches from the waiting to the ready state (for example, at completion of I/O)
 - 4. terminates
- ▶ Nonpreemptive or cooperative scheduling scheme:
 - Process keeps the CPU until released
 - Scheduling for 1 and 4 only (no choice in terms of scheduling)
- ▶ Preemptive: switching states by scheduling scheme (for 2 and 3). Modern OS.
 - results in race conditions when data are shared among several processes (synchronization: ch.6)
 - Ex. Process A updates the data → preempted → process B reads the data (before updated)

5.1 Basic Concepts

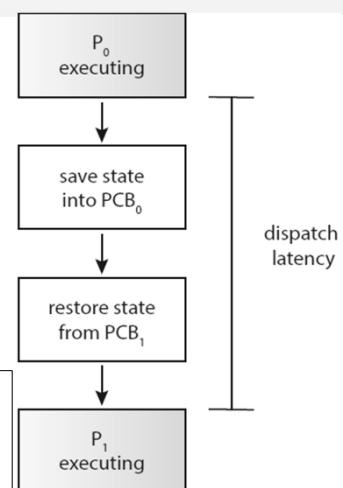
● Dispatcher

- ▶ gives control of the CPU's core to the process
- ▶ functions
 - Switching context
 - Switching to user mode
 - Jumping to the proper location in the user program to resume
- ▶ dispatch latency: time to stop one process and start another
 - should be as short as possible
- ▶ Frequency of the context switches
 - Linux: vmstat 1 3 : 3 lines of output over a 1-second delay

```
-----cpu-----
24      : average number of context switches over 1 second since the system booted
225     : number of context switches in the past second
339     : number of context switches in the second prior to above
```

- cat /proc/2166/status : list various statistics for the process with pid = 2166

```
voluntary ctxt switches      150
nonvoluntary ctxt switches   8      : preempted
```



5.2 Scheduling Criteria

- Affects the choice of scheduling algorithm
- CPU utilization
 - ▶ Rate of CPU being busy
 - ▶ Ideally, 0 to 100 percent. In a real system, from 40 percent (for a lightly loaded system) to 90 percent (for a heavily loaded system).
 - ▶ Linux, macOS, and UNIX: `top` command for CPU utilization
- Throughput
 - ▶ Number of completed processes per time unit
- Turnaround time
 - ▶ Time to execute: from submission to completion
 - ▶ sum of the periods for waiting in the ready queue, executing on the CPU, and doing I/O.
- Waiting time
 - ▶ sum of the periods for waiting in the ready queue
 - ▶ affected by scheduling algorithm



5.2 Scheduling Criteria

- Response time.
 - ▶ Time to start responding, not the time to output the response.
 - ▶ Interactive system
- It is desirable
 - ▶ to maximize CPU utilization and throughput
 - ▶ to minimize turnaround time, waiting time, and response time
- Optimization
 - ▶ Average measure in most cases, or
 - ▶ minimum or maximum values
 - Ex. To guarantee service time, to minimize the maximum response time
 - ▶ for interactive systems (such as a PC desktop or laptop system), it is more important to minimize the variance in the response time than to minimize the average response time
 - Little work has been done



5.3 Scheduling Algorithms

● Assumption

- ▶ one CPU burst (in milliseconds) per process
- ▶ comparing the average waiting time
- ▶ only one processing core available

● First-Come, First-Served (FCFS) Scheduling

- ▶ Simplest
- ▶ FIFO queue
- ▶ convoy effect:
 - All the other processes wait for the one big process to get off the CPU
 - Ex. many I/O-bound processes and one CPU-bound process
 - Lower CPU and device utilization
 - Shorter processes go first
- ▶ Nonpreemptive
 - troublesome for interactive systems

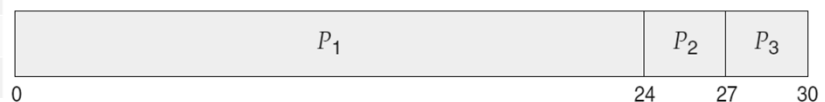
5.3 Scheduling Algorithms

● First-Come, First-Served (FCFS) Scheduling

▶ Ex.

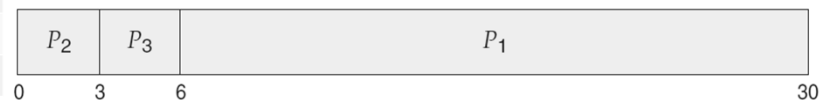
Process	Burst Time
P ₁	24
P ₂	3
P ₃	3

average waiting time: 17



Process	Burst Time
P ₂	3
P ₃	3
P ₁	24

average waiting time: 3



5.3 Scheduling Algorithms

● Shortest-Job-First (SJF) Scheduling

▶ *shortest-next-CPU-burst*

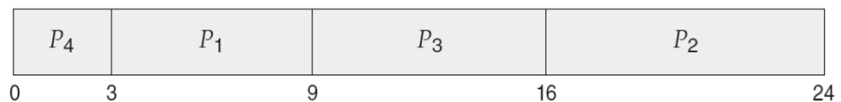
▶ Dispatch the process with the smallest next CPU burst; FCFS for tie break

▶ gives the minimum average waiting time

▶ Nonpreemptive or preemptive (shortest-remaining-time-first: SRTF)

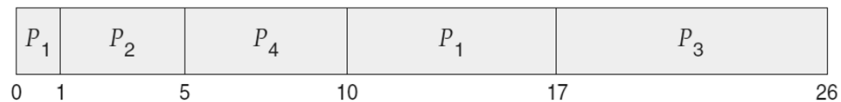
Process	Burst Time
P ₁	6
P ₂	8
P ₃	7
P ₄	3

average waiting time(SJF): 7
average waiting time(FCFS): 10.25



Process	Arrival Time	Burst Time
P ₁	0	8
P ₂	1	4
P ₃	2	9
P ₄	3	5

average waiting time(SJF): 7.75
average waiting time(SRTF): 6.5



5.3 Scheduling Algorithms

● Shortest-Job-First (SJF) Scheduling

▶ Implementation problem: no way to know the length of the next CPU burst

- approximate SJF scheduling: the next CPU burst \approx the previous ones
- exponential average of the measured lengths of previous CPU bursts

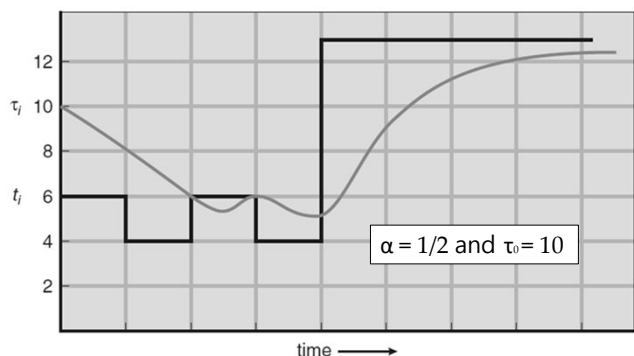
$$-\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$$

– τ_{n+1} : predicted value for the next CPU burst

– t_n : length of the n^{th} CPU burst. most recent

– α : $0 \leq \alpha \leq 1$

– τ_n : stores the past history



CPU burst (t_i)	6	4	6	4	13	13	13	...
"guess" (τ_i)	10	8	6	6	5	9	11	12

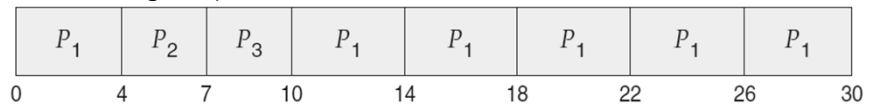
5.3 Scheduling Algorithms

● Round-Robin Scheduling

- ▶ Preemptive: to enable the system to switch between processes
- ▶ Time quantum (or time slice): generally, 10 ~ 100 ms.
- ▶ Ready queue: FIFO with circular operation.
- ▶ The CPU scheduler allocates the CPU for up to 1 time quantum using a timer
 - CPU burst < 1 TQ: the process releases CPU (T or W)
 - CPU burst > 1 TQ: context switch by interrupt (timer) (R)

Process	Burst Time
P ₁	24
P ₂	3
P ₃	3

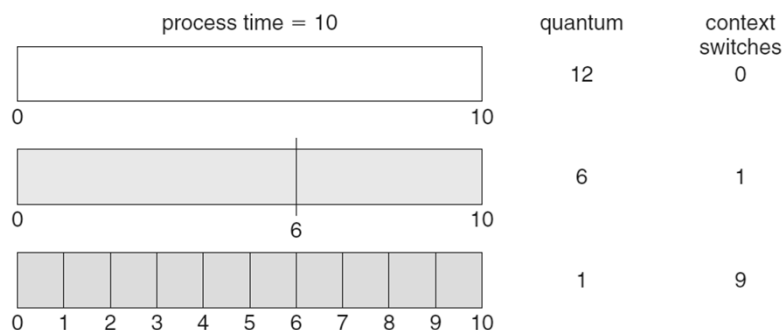
average waiting time: 5.66
small average response time



5.3 Scheduling Algorithms

● Round-Robin Scheduling

- ▶ The size of the time quantum
 - Not too small, not too large: rule of thumb: 80% of CPU bursts < TQ
 - Context switching time: typically, < 10us (< 0.1% of TQ)



- Turnaround time
 - Ex. 3 processes, T=10, TQ=1: average turnaround time=29
 - TQ=10: average turnaround time=20 (FCFS)

5.3 Scheduling Algorithms

● Priority Scheduling

- ▶ Priority: generally indicated by fixed range of numbers.
 - Assigned to each process
 - High: may be '0' or max. number
- ▶ CPU is allocated to the process with the highest priority
 - FCFS: equal priorities for all processes
 - SJF: priority (p) is the inverse of the (predicted) next CPU burst
- ▶ Ex.

Process	Burst Time	Priority
P ₁	10	3
P ₂	1	1
P ₃	2	4
P ₄	1	5
P ₅	5	2

average waiting time: 8.2



5.3 Scheduling Algorithms

● Priority Scheduling

- ▶ Preemptive or nonpreemptive
- ▶ Starvation (indefinite blocking) of low priority process
 - Aging: gradually increasing the priority of long-waiting processes
- ▶ To combine round-robin and priority scheduling
 - runs processes with the same priority using round-robin scheduling
 - Ex.

Process	Burst Time	Priority
P ₁	4	3
P ₂	5	2
P ₃	8	2
P ₄	7	1
P ₅	3	3

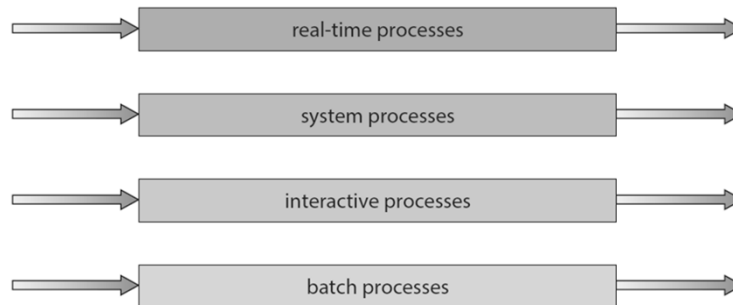


5.3 Scheduling Algorithms

● Multilevel Queue Scheduling

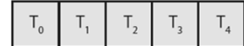
- ▶ Multilevel scheduling using multiple queues
- ▶ Scheduling among the queues and within each queue
- ▶ Ex. Separate queues for each priority
round-robin within a queue ✓
- ▶ separate queues based on the process type

highest priority



lowest priority

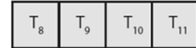
priority = 0



priority = 1



priority = 2



•
•
•

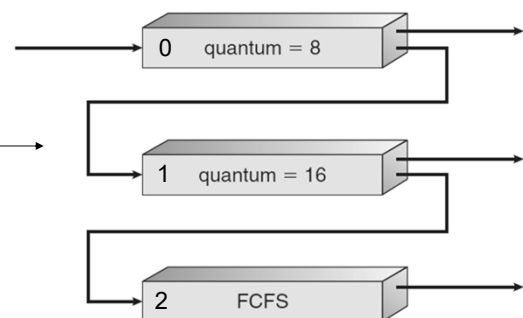
priority = n



5.3 Scheduling Algorithms

● Multilevel Feedback Queue Scheduling

- ▶ a process moves between queues
- ▶ The most general CPU-scheduling algorithm
 - most complex
- ▶ Ex. Long CPU-burst process to lower-priority queue
long-waiting process to higher-priority queue
- ▶ EX. →
- ▶ Parameters
 - # of queues
 - scheduling algorithm for each queue
 - when to upgrade to a higher priority queue
 - when to demote to a lower priority queue
 - starting queue



5.4 Thread Scheduling

- Scheduling issues involving user-level and kernel-level threads
- Contention Scope
 - ▶ Process-contention scope (PCS): scheduling by thread library for user level threads to an LWP
 - many-to-one or many-to-many
 - Typically, based on priority set by programmers
 - ▶ System-contention scope (SCS): scheduling by OS LWP's kernel thread onto a CPU core
- Pthread Scheduling
 - ▶ PTHREAD_SCOPE_PROCESS: PCS scheduling.
 - ▶ PTHREAD_SCOPE_SYSTEM: SCS scheduling.

5.4 Thread Scheduling

```
#include <pthread.h>
#include <stdio.h>
#define NUM_THREADS 5
int main(int argc, char *argv[])
{
    int i, scope;
    pthread_t tid[NUM_THREADS];
    pthread_attr_t attr;
    /* get the default attributes */
    pthread_attr_init(&attr);
    /* first inquire on the current scope */
    if (pthread_attr_getscope(&attr, &scope) != 0)
        fprintf(stderr, "Unable to get scheduling scope\n");
    else {
        if (scope == PTHREAD_SCOPE_PROCESS)
            printf("PTHREAD_SCOPE_PROCESS");
        else if (scope == PTHREAD_SCOPE_SYSTEM)
            printf("PTHREAD_SCOPE_SYSTEM");
        else
            fprintf(stderr, "Illegal scope value.\n");
    }
}

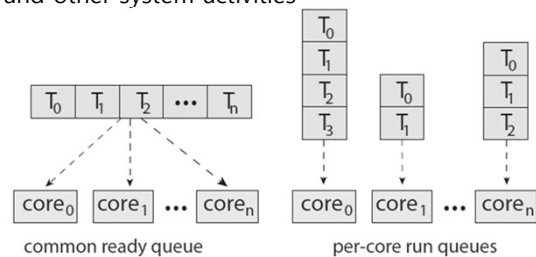
/* set the scheduling algorithm to PCS or SCS */
pthread_attr_setscope(&attr, PTHREAD_SCOPE_SYSTEM);
/* create the threads */
for (i = 0; i < NUM_THREADS; i++)
    pthread_create(&tid[i], &attr, runner, NULL);
/* now join on each thread */
for (i = 0; i < NUM_THREADS; i++)
    pthread_join(tid[i], NULL);
}

/* Each thread will begin control in this function */
void *runner(void *param)
{
    /* do some work ... */
    pthread_exit(0);
}
```

- Linux and macOS systems allow only PTHREAD_SCOPE_SYSTEM

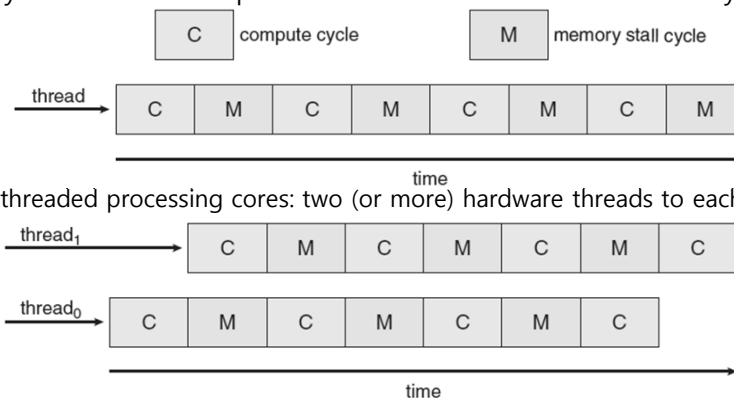
5.5 Multi-Processor Scheduling

- Multiprocessor: scheduling issues become more complex
 - ▶ Multicore CPUs
 - ▶ Multithreaded cores
 - ▶ NUMA systems
 - ▶ Heterogeneous multiprocessing
- Approaches to Multiple-Processor Scheduling
 - ▶ asymmetric multiprocessing
 - Master server: all scheduling decisions, I/O processing, and other system activities
 - Others: Execute users codes
 - potential bottleneck by the master server
 - ▶ Symmetric multiprocessing (SMP)
 - each processor is self-scheduling.
 - possible strategies
 1. All threads may be in a common ready queue.
 - possible race condition
 2. Each processor may have its own private queue of threads.
 - load balancing algorithm



5.5 Multi-Processor Scheduling

- Multicore Processors
 - ▶ SMP systems with multicore processors: faster and consuming less power
 - ▶ memory stall: cache miss: processors are much faster than memory

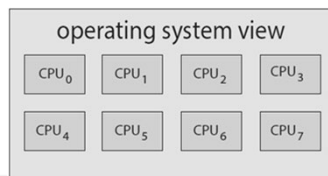
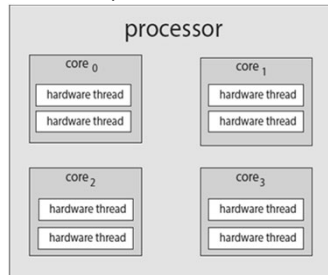


5.5 Multi-Processor Scheduling

● Multicore Processors

▶ chip multithreading (CMT)

- Intel: hyper-threading (aka simultaneous multithreading or SMT): 2 threads per core
- Oracle Sparc M7 (8 cores): 8 threads per core → 64 logical cores



AMD Ryzen™ 9 5950X

CPU 코어 수: 16

스레드 수: 32

최대 부스트 클럭: 최대 4.9GHz

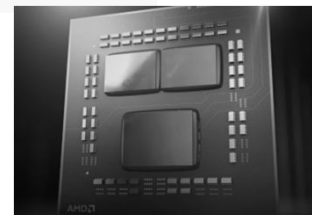
기본 클럭: 3.4GHz

총 L2 캐시: 8MB

총 L3 캐시: 64MB

기본 TDP/TDP: 105W

Processor Technology for CPU Cores: TSMC 7nm FinFET



Soongsil Univ.
Digital System Design Lab.

Operating Systems

Copyright 2023. (차세대반도체 혁신공유대학 사업단) all rights reserved 23

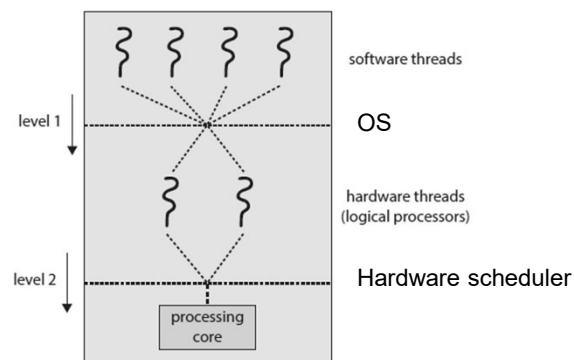
23

5.5 Multi-Processor Scheduling

● Multicore Processors

▶ Multithreading a processing core

1. Coarse grained: execution until a long-latency event.
 - Pipeline flushing → high cost of thread switching
 2. Fine-grained (interleaved): thread switching at the boundary of instruction cycle
 - No pipeline flushing → small cost; extra logic required
- Concurrent, not parallel:
 - Holding multiple threads;
 - sharing hardware for execution
 - Two levels of scheduling
 - sharing-aware algorithm
 - » Ex. 2 cores and 2 threads/core



Soongsil Univ.
Digital System Design Lab.

Operating Systems

Copyright 2023. (차세대반도체 혁신공유대학 사업단) all rights reserved 24

24

5.5 Multi-Processor Scheduling

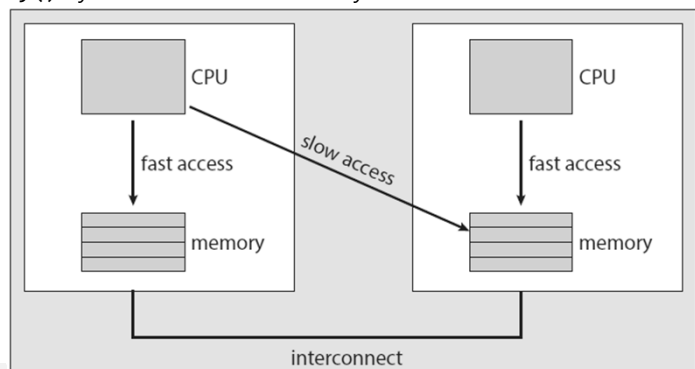
● Load Balancing

- ▶ evenly distribute workload across all processors
- ▶ Approaches
 1. Push migration: periodical checking → moving (or pushing) threads if necessary
 2. Pull migration: an idle processor pulls a waiting task
 - not mutually exclusive: ex. Linux CFS scheduler, FreeBSD ULE scheduler
- ▶ Balanced load
 - approximately the same number of threads in all queues or
 - equal distribution of thread priorities across all queues

5.5 Multi-Processor Scheduling

● Processor Affinity

- ▶ to keep a thread running on the same processor and take advantage of a warm cache
- ▶ Easy for private, per-processor ready queues
- ▶ Processor affinity forms
 1. Soft affinity: process running on the same processor is not guaranteed
 2. Hard affinity: a process specifies a subset of processors using system calls
 - Linux: soft affinity and `sched_setaffinity()` system call for hard affinity
- ▶ non-uniform memory access (NUMA)
 - NUMA-aware algorithm
 - Process scheduling to near CPU
- ▶ Counteraction between load balancing and processor affinity



5.5 Multi-Processor Scheduling

● Heterogeneous Multiprocessing (HMP)

- ▶ ARM's big.LITTLE architecture and Intel's hybrid technology (P/E core)
 - Qualcomm SD8Gen2: 1 x Cortex-X3 + 2 x Cortex-A715 + 2 x Cortex-A710 + 3 x Cortex-A510
 - Intel i9 13900k: 8P + 16E
- ▶ Power management
- ▶ Big/P cores: interactive tasks, high performance and short burst tasks
 - disabled in power-saving mode
- ▶ LITTLE/E cores: low performance or long bursts task, background jobs
- ▶ Windows 10 supports HMP scheduling
 - allowing a thread to select a scheduling policy that best supports its power management demands

Summary

- CPU scheduling
 - ▶ selecting a waiting process from the ready queue and allocating the CPU using the dispatcher.
- Scheduling algorithms:
 - ▶ either preemptive or nonpreemptive
 - ▶ Evaluating criteria
 - (1) CPU utilization, (2) throughput, (3) turnaround time, (4) waiting time, and (5) response time.
 - ▶ FCFS: simplest; not efficient for waiting time
 - ▶ SJF: optimal for WT; difficult in implementing
 - ▶ RR: preemption, time quantum, short response time
 - ▶ Priority: the highest priority first.
 - ▶ Multilevel queue: multiple queues by priority, different scheduling in each queue
 - ▶ Multilevel feedback queues: process migration between queues
- Multicore processors with multiple hardware threads (logical CPUs)
 - ▶ Load balancing and processor affinity

Exercises, problems and projects

● Exercises

▶ 5.4, 5.6, 5.7, 5.8

● Problems

▶ 5.13, 5.21

5.13 One technique for implementing lottery scheduling works by assigning processes lottery tickets, which are used for allocating CPU time. Whenever a scheduling decision has to be made, a lottery ticket is chosen at random, and the process holding that ticket gets the CPU. The BTV operating system implements lottery scheduling by holding a lottery 50 times each second, with each lottery winner getting 20 milliseconds of CPU time ($20 \text{ milliseconds} \times 50 = 1 \text{ second}$). Describe how the BTV scheduler can ensure that higher-priority threads receive more attention from the CPU than lower-priority threads.

5.21 Consider a variant of the RR scheduling algorithm in which the entries in the ready queue are pointers to the PCBs.

- What would be the effect of putting two pointers to the same process in the ready queue?
- What would be two major advantages and two disadvantages of this scheme?
- How would you modify the basic RR algorithm to achieve the same effect without the duplicate pointers?

