
 Soongsil Univ.
<http://design.ssu.ac.kr>

 Digital
System
Design Lab.

Ch. 10 Virtual Memory


Chanho Lee
Soongsil University

Copyright 2023. (차세대반도체 혁신공유대학 사업단) all rights reserved

1

Objectives

- Virtual memory and benefits.
- Page loading on demand
- Page-replacement algorithms.
- Working set of a process and program locality.
- Virtual memory in Linux, Windows 10

 Soongsil Univ.
Digital System Design Lab.

Operating Systems

Copyright 2023. (차세대반도체 혁신공유대학 사업단) all rights reserved 2

2

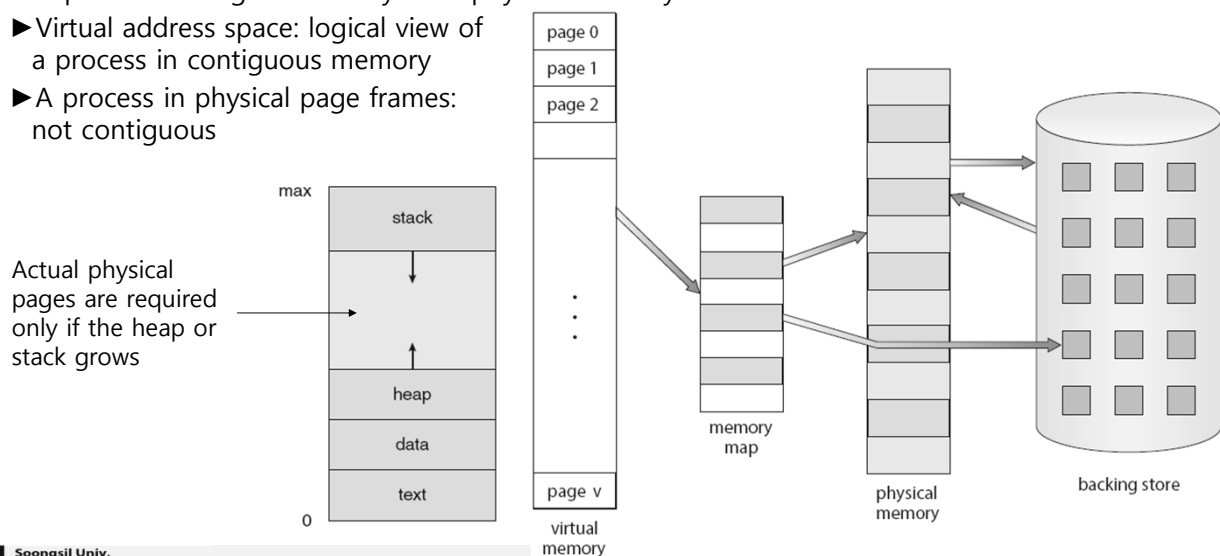
10.1 Background

- The instructions being executed must be in physical memory
 - ▶ limits the size of a program to the size of physical me
 - ▶ the entire program is not needed in many cases
 - code to handle unusual error conditions: rare error occurrence → almost never executed
 - Arrays, lists, and tables: often allocated more for the worst case.
 - Certain options and features of a program may be used rarely.
 - Frequently used functions are not used at the same time
 - ▶ Partial loading of a program to memory
 - Programming an extremely large virtual address space → simplifying the programming task
 - Using less physical memory → increase the degree of multiprogramming, CPU utilization and throughput
 - Less I/O for loading or swapping → programs runs faster

3

10.1 Background


- Virtual memory
 - ▶ Separation of logical memory from physical memory.
 - ▶ Virtual address space: logical view of a process in contiguous memory
 - ▶ A process in physical page frames: not contiguous



4

10.1 Background


- Sparse address spaces: virtual address spaces that include holes
 - ▶ beneficial because
 - the holes can be filled as the stack or heap segments grow or
 - if we wish to dynamically link libraries (or possibly other shared objects) during program execution
- Page sharing
 - ▶ files and memory to be shared
 - system libraries
 - sharing memory (IPC)
 - ▶ Part of virtual address space, but in the shared physical memory




5

10.2 Demand Paging


- Pages are loaded only when demanded
- Basic Concepts
 - ▶ Pages in memory and in secondary storage
 - Need some hardware support to distinguish the two
 - Valid: both legal and in memory
 - Invalid: either is not valid or in secondary storage
 - no effect until accessing the page
 - Access results in a page fault




logical memory



page table



physical memory

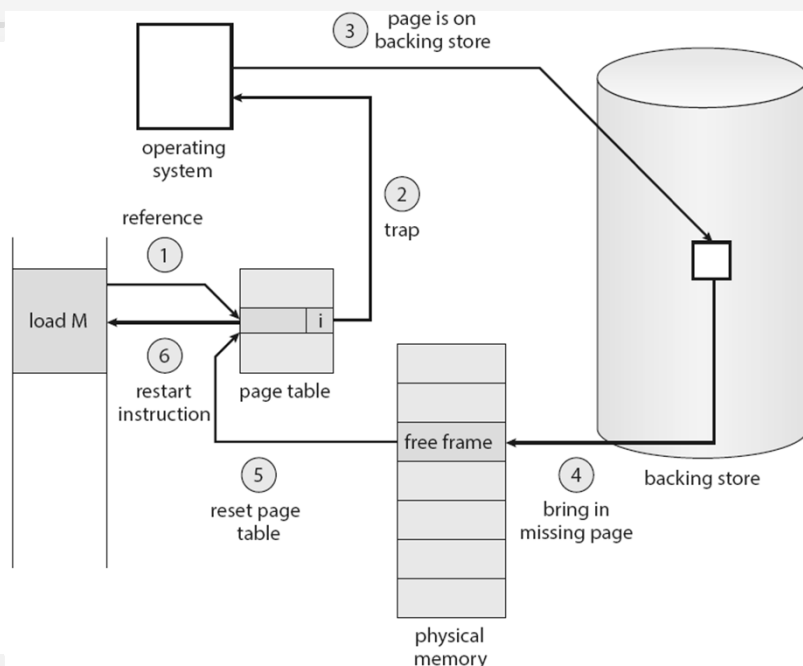


backing store

6

10.2 Demand Paging

- ▶ ② Handling page fault
 - check valid or invalid memory access using an internal table
 - Invalid: termination
 - Valid: page-in ③



10.2 Demand Paging

- Pure demand paging
 - ▶ Process start with no pages in memory → page faults on the first instruction
 - one or more page faults for instruction and data
 - locality of reference
- Hardware to support demand paging
 - ▶ Page table and secondary memory
 - ▶ Ability to restart any instruction after a page fault
- Free-Frame List
 - ▶ a pool of free frames $\text{head} \rightarrow 7 \rightarrow 97 \rightarrow 15 \rightarrow 126 \dots \rightarrow 75$
 - ▶ zero-fill-on-demand: "zeroed-out" before being allocated
- Performance of Demand Paging
 - ▶ effective access time = $(1 - p) \times ma + p \times \text{page fault time}$
 - *ma*: memory-access time (200ns)
 - *p*: probability of a page fault ($0 \leq p \leq 1$)

10.2 Demand Paging


- ▶ page fault resolving sequence (handler): refer to the textbook for the sequences in detail
 - 1. Service the page-fault interrupt (1 ~ 100us)
 - 2. Read in the page (~8ms for HDD)
 - 3. Restart the process (1 ~ 100us)
- ▶ effective access time = $(1 - p) \times (200) + p (8 \text{ ms})$

$$= (1 - p) \times 200 + p \times 8,000,000$$

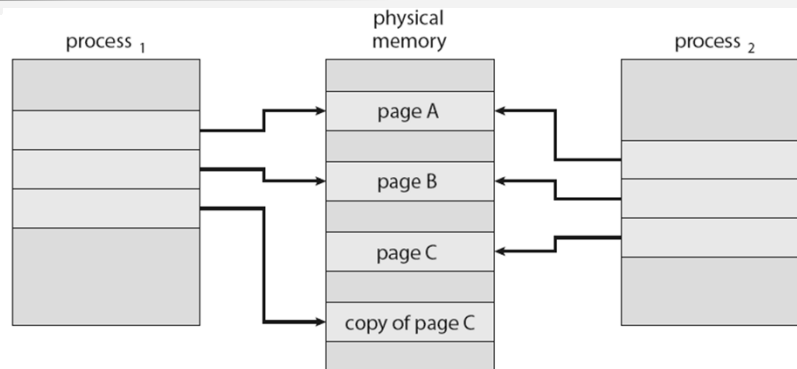
$$= 200 + 7,999,800 \times p.$$
 - directly proportional to the page-fault rate
 - If $p=0.001$, 8.2us \rightarrow performance degradation by a factor of 40
 - <10% degradation $\rightarrow p < 0.0000025$
- ▶ Swap space (generally faster than the file system)
 - copying an entire file image into the swap space at process startup
 - demand-page from the file system initially but write the pages to swap space as they are replaced
 - Windows, Linux

10.3 Copy-on-Write

- fork() system call: page sharing \rightarrow bypass demand paging
 - ▶ rapid process creation and minimizes the number of page allocation
 - ▶ fork()-exec(): copying of the parent's address space may be unnecessary
- copy-on-write:
 - ▶ parent and child processes initially to share the same pages.
 - ▶ shared pages are marked as copy-on-write \rightarrow if either process writes to it, create a copy



10.3 Copy-on-Write




- ▶ only the pages that are modified by either process are copied
- ▶ Windows, Linux, and macOS

11

10.4 Page Replacement

- ▶ demand paging → increase degree of multiprogramming → over-allocating memory
- ▶ System memory: program pages + I/O buffer → out of free frames
- ▶ swapping pages with page replacement






12

10.4 Page Replacement

● Basic Page Replacement

- ▶ Replacement: requires two page transfers
- ▶ modify bit (or dirty bit)
 - Not set: no page out
- ▶ frame-allocation algorithm
 - No. of frames to allocate to process
- ▶ page-replacement algorithm
 - select the frames to be replaced




13

10.4 Page Replacement

● FIFO Page Replacement

- ▶ replace the page at the head of the queue



reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2	2	2	4	4	4	0	0	0	0	0	0	0	7	7	7
	0	0	0	3	3	3	2	2	2	1	1	1	1	1	1	1	0	0
		1	1	1	0	0	0	0	3	3	2	2	2	2	2	2	2	1

page frames

15 page faults

14


10.4 Page Replacement

● Optimal Page Replacement

- ▶ To overcome Belady's anomaly
- ▶ Replace the page not be used for the longest time
- ▶ the lowest possible page fault rate for a fixed number of frames

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



9 page faults

page frames

- ▶ difficult to implement: it requires future knowledge of the reference string

10.4 Page Replacement

● Least recently used (LRU) Page Replacement

- ▶ approximation of the optimal algorithm
- ▶ use the recent past as an approximation of the near future
→ the page that has not been used for the longest time

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1




page frames

12 page faults

reference string

4 7 0 7 1 0 1 2 1 2 7 1 2



stack
before
a

stack
after
b


- ▶ Implementing may require substantial hardware assistance
 - Counters: store the "time" of the last reference to each page → a page with the smallest is replaced
 - Stack: the most recently used page at the top of the stack → a page at the bottom is replaced

10.5 Allocation of Frames

- Example: simple strategy
 - ▶ 128 frames, 35 for kernel, 93 for user, pure demanding page: allocation?
 - ▶ 93 page faults, page replacements, termination, 93 free frames
- Minimum Number of Frames
 - ▶ No. of frames allocated $\downarrow \rightarrow$ page-fault rate $\uparrow \rightarrow$ performance \downarrow
 - ▶ Allocate frames to hold all the pages that any single instruction can reference
 - defined by the processor architecture
- Allocation Algorithms
 - ▶ Equal allocation: m/n frames (split m frames among n processes)
 - ▶ Proportional allocation: allocate available memory to each process according to its size.
 - Priority can be combined
- Global versus Local Allocation
 - ▶ From the set of all frames
 - ▶ From its own set of allocated frames

10.6 Thrashing

- Thrashing
 - ▶ High paging activity: spending more time paging than executing
 - \rightarrow severe performance problems
 - ▶ Ex. replace a page that will be needed again right away \rightarrow quickly faults again, and again
- Cause of Thrashing
 - ▶ Low CPU utilization \rightarrow increase the degree of multiprogramming by introducing a new process
 - ▶ With a global page-replacement algorithm, new process takes frames \rightarrow more processes fault \rightarrow waiting for paging device \rightarrow ready queue empties \rightarrow low CPU utilization
 - ▶ Local replacement algorithm (or priority replacement algorithm)
 - limits the effects of thrashing
 - Thrashing process is mainly in the waiting queue
 - The effective access time for paging device will increase



10.7 Memory Compression

- Compress several frames into a single frame, enabling the system to reduce memory usage → No swapping

▶ Before compression

free-frame list

head → 7 → 2 → 9 → 21 → 27 → 16

modified frame list

head → 15 → 3 → 35 → 26

▶ After compression

free-frame list

head → 2 → 9 → 21 → 27 → 16 → 15 → 3 → 35

modified frame list

head → 26

compressed frame list

head → 7

10.9 Other Considerations

- Prepaging
 - ▶ to bring some—or all—of the pages that will be needed at one time ↔ pure demand paging
 - ▶ to prevent this high level of initial paging
- Page Size
 - ▶ Memory is better utilized with smaller pages ↔ increases the size of the page table (per process)
 - Better resolution for locality:
 - ▶ Time for read/write pages
 - HDD: seek and latency times are much larger than transfer time (larger page size is desirable)
 - SSD: no seek time and much smaller latency. Different read and write time.
 - ▶ TLB Reach (related to hit ratio)
 - The number of entries multiplied by the page size: memory size accessible from the TLB
 - Larger page size is desirable

10.9 Other Considerations

● Program Structure

- ▶ Page size: 128 words
- ▶ Row major: 128 x 128 page faults

```
int i, j;
int[128][128] data;
for (j = 0; j < 128; j++)
    for (i = 0; i < 128; i++)
        data[i][j] = 0; // zeros one word in each page
```

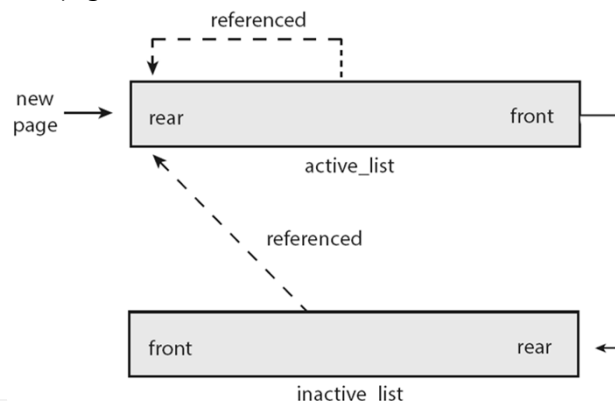
- ▶ Column major: 128 page faults

```
int i, j;
int[128][128] data;
for (i = 0; i < 128; i++)
    for (j = 0; j < 128; j++)
        data[i][j] = 0; // zeros all the word in a page
```

10.10 Operating-System Examples

● Linux

- ▶ demand paging, allocating pages from a list of free frames
- ▶ global page-replacement policy similar to the LRU-approximation clock algorithm
- ▶ page lists: an active list and an inactive list
 - inactive list: contains pages that have not recently been referenced and are eligible to be reclaimed
 - Each page has an accessed bit: set whenever the page is referenced
- ▶ page-out daemon process kswapd
 - periodically checks the amount of free memory
 - Free memory falls below a certain threshold
 - scanning pages in the inactive list and reclaiming them for the free list



10.10 Operating-System Examples

● Windows

- ▶ Windows 10 supports 32- and 64-bit systems running on Intel (IA-32 and x86-64) and ARM
- ▶ 32-bit systems
 - the default virtual address space: 2 GB,
 - can be extended to 3 GB.
 - 4 GB of physical memory
- ▶ 64-bit systems:
 - 128-TB virtual address space
 - 24 TB of physical memory (Server: 128 TB)
- ▶ shared libraries, demand paging, copy-on-write, paging, and memory compression
- ▶ demand paging with clustering
 - Bringing faulting page + several pages immediately preceding and following the faulting page
 - Data page: cluster=3; others: cluster=7

Summary

- Virtual memory
 - ▶ Definition and benefits
- Demand paging
 - ▶ A page fault
- Copy-on-write
- page-replacement algorithm
 - ▶ FIFO, optimal, and LRU. Pure LRU, LRU-approximation algorithms.
 - ▶ Global page-replacement algorithms vs. local page-replacement algorithms
- Thrashing
- Memory compression
- Other Considerations
 - ▶ Prepaging, page size, TLB reach, Program structure
- OS examples: Linux, Windows

Exercises, problems and projects

● Exercises

▶ 10.1, 10.5, 10.7

● Problems

▶ 10.16

10.16 A simplified view of thread states is ready, running, and blocked, where a thread is either ready and waiting to be scheduled, is running on the processor, or is blocked (for example, waiting for I/O).

Assuming a thread is in the running state, answer the following questions, and explain your answers:

- Will the thread change state if it incurs a page fault? If so, to what state will it change?
- Will the thread change state if it generates a TLB miss that is resolved in the page table? If so, to what state will it change?
- Will the thread change state if an address reference is resolved in the TLB? If so, to what state will it change?

